
Finding Dense Subgraphs via Low-Rank Bilinear Optimization

Dimitris S. Papailiopoulos

Ioannis Mitliagkas

Alexandros G. Dimakis

Constantine Caramanis

The University of Texas at Austin

DIMITRIS@UTEXAS.EDU

IOANNIS@UTEXAS.EDU

DIMAKIS@AUSTIN.UTEXAS.EDU

CONSTANTINE@UTEXAS.EDU

Abstract

Given a graph, the Densest k -Subgraph (DkS) problem asks for the subgraph on k vertices that contains the largest number of edges. In this work, we develop a new algorithm for DkS that searches a low-dimensional space for provably good solutions. Our algorithm comes with novel performance bounds that depend on the graph spectrum. Our graph-dependent bounds are in practice significantly tighter than worst case *a priori* bounds: for most tested real-world graphs we find subgraphs with density provably within 70% of the optimum.

Our algorithm runs in nearly linear time, under spectral assumptions satisfied by most graphs found in applications. Moreover, it is highly scalable and parallelizable. We demonstrate this by implementing it in MapReduce and executing numerous experiments on massive real-world graphs that have up to billions of edges. We empirically show that our algorithm can find subgraphs of significantly higher density compared to the previous state of the art.

1. Introduction

Given a graph \mathcal{G} on n vertices with m edges and a parameter k , we are interested in finding an induced subgraph on k vertices with the largest average degree, also known as the *maximum density*. This is the *Densest k -Subgraph* (DkS) – a fundamental problem in combinatorial optimization with applications in numerous fields including social sciences, communication networks, and biology (see *e.g.* (Hu et al., 2005; Gibson et al., 2005; Dourisboure et al., 2007; Saha et al., 2010; Miller et al., 2010; Bahmani et al., 2012)).

Proceedings of the 31st International Conference on Machine Learning, Beijing, China, 2014. JMLR: W&CP volume 32. Copyright 2014 by the author(s).

DkS is a notoriously hard problem. It is NP-hard by reduction to MAXCLIQUE. Moreover, Khot showed in (Khot, 2004) that, under widely believed complexity theoretic assumptions, DkS cannot be approximated within an arbitrary constant factor.¹ The best known approximation ratio was $n^{1/3+\epsilon}$ (for some small ϵ) due to (Feige et al., 2001). Recently, (Bhaskara et al., 2010) introduced an algorithm with approximation ratio $n^{1/4+\epsilon}$, that runs in time $n^{O(1/\epsilon)}$. Such results, where the approximation factor scales as a polynomial in the number of vertices, are too pessimistic for real-world applications. The resistance to better approximation despite the long history of the problem suggests that DkS is probably very hard in the worst case.

Our Contributions. In this work we move beyond the worst case framework. We present a novel DkS algorithm that has two key features: *i*) it comes with approximation guarantees that are surprisingly tight on real-world graphs and *ii*) it is fully parallelizable and can scale up to graphs with billions of edges.

Our algorithm combines spectral and combinatorial techniques; it relies on examining candidate subgraphs obtained from vectors lying in a low-dimensional subspace of the adjacency matrix of the graph. This is accomplished through a framework called the *Spanogram*, which we define below. At its core our algorithm has an exact solver for constant rank *bilinear* programs with combinatorial constraints.

Our approximation guarantees are *data-dependent*: they are related to the spectrum of the adjacency matrix. Let opt denote the average degree (i.e., the density) of the densest k -subgraph, where $0 \leq \text{opt} \leq k-1$. Our algorithm takes as input the graph, the subgraph size k and an accuracy parameter $d \in \{1, \dots, n\}$ and runs in time $O(kn^{d+1})$. Note that under a spectral condition satisfied by many real graphs, the computational time is reduced to nearly-linear as we subse-

¹approximation ratio ρ means that there exists an algorithm that produces in polynomial time a number A , such that $1 \leq \frac{\text{opt}}{A} \leq \rho$, where opt is the optimal density.

quently discuss. The output is a subgraph on k vertices with density opt_d , for which we obtain the following approximation result:

Theorem 1. *For any unweighted graph*

$$\text{opt}_d \geq \text{opt}/(2 + o_n(1)) - 2 \cdot |\lambda_{d+1}|.$$

where λ_i is the i th largest, in magnitude, eigenvalue of the adjacency matrix of the graph. Moreover, if the graph is bipartite, or if the largest d eigenvalues of the graph are positive, then $\text{opt}_d \geq \text{opt} - 2 \cdot |\lambda_{d+1}|$.

Our bounds come close to $2 + \epsilon$ and $1 + \epsilon$ factor approximations, when λ_{d+1} is significantly smaller than the density of the densest k -subgraph. In the following theorem, we give such an example. However, we would like to note that in the worst case our bounds might not yield something meaningful.

Theorem 2. *If the densest- k -subgraph contains a constant fraction of all the edges, and $k = \Theta(\sqrt{E})$, then we can approximate DkS within a factor of $2 + \epsilon$, in time $n^{O(1/\epsilon^2)}$. If additionally the graph is bipartite, we can approximate DkS within a factor of $1 + \epsilon$.*

The above result is similar to the $1 + \epsilon$ approximation ratio of (Arora et al., 1995) for dense graphs, where the densest- k -subgraph contains a constant fraction of the $\Omega(n^2)$ edges, where $k = \Omega(n)$. The innovation here is that our ratio also applies to *sparse graphs* with sublinear number of edges.

Computable upper bounds. In addition to these theoretical guarantees, our analysis allows us to obtain a graph-dependent upper bound for the optimal subgraph density. This is shown in Fig. 3 in our experimental section, where for many graphs our algorithm is provably within 70% from the upper bound of opt . These are far stronger guarantees than the best available *a priori* bounds. This illustrates the potential power of graph-dependent guarantees that, however, require the execution of an algorithm.

Nearly-linear time approximation. Our algorithm has a worst-case running time of $O(kn^{d+1})$, where d controls the quality of the approximation. However, under some mild spectral assumptions, a randomized version of our algorithm runs in nearly-linear time.

Theorem 3. *Let the d largest eigenvalues of the graph be positive and let the d and $d+1$ largest have constant ratio: $\lambda_d/\lambda_{d+1} = \Theta(1)$. Then, we can modify our algorithm to output, with high probability, a subgraph with density $(1 - \frac{1}{\log n}) \cdot \text{opt}_d$ in nearly-linear time $O(m \cdot \log n + n \cdot \log^{d+1} n)$, where m is the number of edges.*

We found that the above spectral condition holds for all $d \leq 5$, in many real-world graphs that we tested.

Scalability. We develop two key scalability features that allow us to scale up efficiently on massive graphs.

Vertex sparsification: We introduce a pre-processing step that eliminates vertices that are unlikely to be part of the densest k -subgraph. The elimination is based on the vertices' *weighted leverage scores* (Mahoney & Drineas, 2009; Boutsidis et al., 2009) and admits a provable bound on the introduced error. We empirically found that even with a negligible additional error, the elimination dramatically reduced problem sizes in all tested datasets.

MapReduce implementation: We show that our algorithm is *fully-parallelizable* and tailor it for the MapReduce framework. We use our MapReduce implementation to run experiments on Elastic MapReduce (EMR) on Amazon. In our large-scale experiments, we were able to scale out to thousands of mappers and reducers in parallel over 800 cores, and find large dense subgraphs in graphs with billions of edges.

1.1. Related work

DkS algorithms: One of the few positive results for DkS is a $1 + \epsilon$ approximation for dense graphs where $m = \Omega(n^2)$, and in the linear subgraph setting $k = \Omega(n)$ (Arora et al., 1995). For some values of $m = o(n^2)$ a $2 + \epsilon$ approximation was established by (Suzuki & Tokuyama, 2005). Moreover, for any $k = \Omega(n)$ a constant factor approximation is possible via a greedy approach by (Asahiro et al., 2000), or via semidefinite relaxations by (Srivastav & Wolf, 1998) and (Feige & Langberg, 2001). Recently, (Alon et al., 2013) established new approximation results for graphs with small “ ϵ -rank,” using an approximate DkS solver for low-rank perturbed versions of the adjacency matrix.

There is a vast literature on algorithms for detecting communities and well-connected subgraphs: greedy schemes (Ravi et al., 1994), optimization approaches (Jethava et al., 2012; d'Aspremont et al., 2010; Ames, 2011), and the truncated power method (Yuan & Zhang, 2011). We compare with various of these algorithms in our evaluation section.

The Spannogram framework: We present an exact solver for *bilinear* optimization problems on matrices of constant rank, under $\{0, 1\}$ and sparsity constraints on the variables. Our theory is a generalization of the Spannogram framework, originally introduced in the foundational work of (Karytinos & Liavas, 2010) and further developed in (Asteris et al., 2014; Papailiopoulos et al., 2013), that obtains exact solvers for low-rank *quadratic* optimization problems with combinatorial constraints, such as sparse PCA.

MapReduce algorithms for graphs: The design of

MapReduce algorithms for massive graphs is an active research area as Hadoop becomes one of the standards for storing large data sets. The related work by Bahmani *et al.* (Bahmani et al., 2012) designs a novel MapReduce algorithm for the *densest subgraph* problem. This densest subgraph problem requires finding a subgraph of highest *normalized density* without enforcing a specific subgraph size k . Surprisingly, without a subgraph size restriction, the densest subgraph becomes polynomially solvable and therefore fundamentally different from what we consider in this paper.

2. Proposed Algorithm

The *density* of a subgraph indexed by a vertex set $\mathcal{S} \subseteq \{1, \dots, n\}$ is equal to the average degree of the vertices within \mathcal{S} :

$$\text{den}(\mathcal{S}) = \frac{\mathbf{1}_S^T \mathbf{A} \mathbf{1}_S}{|\mathcal{S}|}$$

where \mathbf{A} is the adjacency matrix ($A_{i,j} = 1$ if (i,j) is an edge, else $A_{i,j} = 0$) and the indicator vector $\mathbf{1}_S$ has 1s in the entries indexed by \mathcal{S} and 0 otherwise. Observe that $\mathbf{1}_S^T \mathbf{A} \mathbf{1}_S = \sum_{i,j \in S} A_{i,j}$ is twice the number of edges in the subgraph with vertices in \mathcal{S} .

For a fixed subgraph size $|\mathcal{S}| = k$, we can express DkS as a quadratic optimization:

$$\text{DkS : } \text{opt} = (1/k) \cdot \max_{|\mathcal{S}|=k} \mathbf{1}_S^T \mathbf{A} \mathbf{1}_S$$

where $|\mathcal{S}| = k$ denotes that the optimization variable is a k -vertex subset of $\{1, \dots, n\}$.

The bipartite version of DkS. We approximate DkS via approximating its bipartite version. This problem can be expressed as a bilinear maximization:

$$\text{DBkS : } \text{opt}^B = (1/k) \cdot \max_{|\mathcal{X}|=k} \max_{|\mathcal{Y}|=k} \mathbf{1}_X^T \mathbf{A} \mathbf{1}_Y.$$

As we see in the following lemma, the two problems are fundamentally related: a good solution for the bipartite version of the problem maps to a “half as good” solution for DkS. The proof can be found in the supplemental material, where we describe how to convert an algorithm for DBkS to one for DkS.

Lemma 1. *A ρ -approximation algorithm for DBkS implies a 2ρ -approximation algorithm for DkS.*

2.1. DBkS through low rank approximations.

At the core of our approximation lies a constant rank solver: we show that DBkS can be solved in polynomial time on constant rank matrices. We solve constant rank instances of DBkS instead of DkS due to an important implication: DkS is NP-hard even for rank-1 matrices, as we show in the supplemental material.

The exact steps of our low-rank DBkS algorithm are given in the pseudo-code tables referred to as Algorithms 1 and 2.² We continue with a high-level description.

Step 1: Obtain $\mathbf{A}_d = \sum_{i=1}^d \lambda_i \mathbf{v}_i \mathbf{v}_i^T$, a rank- d approximation of \mathbf{A} . Here, λ_i is the i -th largest in magnitude eigenvalue and \mathbf{v}_i the corresponding eigenvector.

Step 2: Use \mathbf{A}_d to obtain $O(n^d)$ candidate subgraphs. For any matrix \mathbf{A} we can solve DBkS by exhaustively checking all $\binom{n}{k} \cdot \binom{n}{k}$ pairs $(\mathcal{X}, \mathcal{Y})$ of k -subsets of vertices. Surprisingly, if we want to find the \mathcal{X}, \mathcal{Y} pairs that maximize $\mathbf{1}_X^T \mathbf{A}_d \mathbf{1}_Y$, i.e., the bilinear problem on the rank- d matrix \mathbf{A}_d , then we show that only $O(n^d)$ candidate pairs need to be examined.

In the next section, we derive the constant rank-solver using two key facts. First, for each fixed vertex set \mathcal{Y} , we show that it is easy to find the optimal set \mathcal{X} that maximizes $\mathbf{1}_X^T \mathbf{A}_d \mathbf{1}_Y$. Since this turns out to be easy, then the challenge is to find the number of different vertex sets \mathcal{Y} that we need to check. Do we need to exhaustively check all $\binom{n}{k}$ supports \mathcal{Y} ? We show that this question is equivalent to searching the span of the first d eigenvectors of \mathbf{A} , and collecting in a set \mathcal{S}_d the top- k coordinates of all vectors in that d -dimensional space. By modifying the Spannogram theory of (Karytinos & Liavas, 2010; Asteris et al., 2014; Papailiopoulos et al., 2013), we show how this set has size $O(n^d)$ and can be constructed in time $O(n^{d+1})$. This will imply that DBkS can be solved in time $O(n^{d+1})$ on \mathbf{A}_d .

Computational Complexity. The worst-case time complexity of our algorithm is $O(T_d + n^{d+1})$, where T_d is the time to compute the first d eigenvectors of \mathbf{A} . Under conditions satisfied by many real world graphs, we show that the complexity reduces to *nearly linear* in the size of \mathcal{G} : $O(m \cdot \log n + n \cdot \log^{d+1} n)$.

Algorithm 1 lowrankDBkS(k, d, \mathbf{A})

```

1:  $[\mathbf{V}_d, \mathbf{\Lambda}_d] = \text{EVD}(\mathbf{A}, d)$ 
2:  $\mathcal{S}_d = \text{Spannogram}(k, \mathbf{V}_d)$ 
3:  $\{\mathcal{X}_d, \mathcal{Y}_d\} = \arg \max_{|\mathcal{X}|=k} \max_{\mathcal{Y} \in \mathcal{S}_d} \mathbf{1}_X^T \mathbf{V}_d \mathbf{\Lambda}_d \mathbf{V}_d^T \mathbf{1}_Y$ 
4: Output:  $\{\mathcal{X}_d, \mathcal{Y}_d\}$ 


---


1:  $\text{Spannogram}(k, \mathbf{V}_d)$ 
2:  $\mathcal{S}_d = \{\text{top}_k(\mathbf{v}) : \mathbf{v} \in \text{span}(\mathbf{v}_1, \dots, \mathbf{v}_d)\}$ 
3: Output:  $\mathcal{S}_d$ .

```

Remark 1. In the following we present an exact solver for DBkS on constant rank approximations of \mathbf{A} . Our general algorithm for DkS makes a number of $O(k)$ calls to the DBkS low-rank solver, on slightly different matrices that are obtained by sub-sampling the

²In the pseudocode of Algorithm 1, $\text{top}_k(\mathbf{v})$, denotes the indices of the k largest signed elements of \mathbf{v} .

entries of a low-rank approximation of \mathbf{A} . The details of our general algorithm can be found in the supplemental material.

2.2. Approximation Guarantees

We approximate DBkS by finding a solution to the constant rank problem

$$\max_{|\mathcal{X}|=k} \max_{|\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}}.$$

We output a pair of vertex sets, $\mathcal{X}_d, \mathcal{Y}_d$, which we refer to as the *rank-d optimal solution*, that has density

$$\text{opt}_d^B = (1/k) \cdot \mathbf{1}_{\mathcal{X}_d}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}_d}.$$

Our approximation guarantees measure how far opt_d^B is from opt^B , the optimal density for DBkS. Our bounds capture a simple core idea: the loss in our approximation comes due to solving the problem on \mathbf{A}_d instead of solving it on the full rank matrix \mathbf{A} . This loss is quantified in the next lemma. The proofs of the following results are in the supplemental material.

Lemma 2. *For any matrix \mathbf{A}*

$$\text{opt}_d^B \geq \text{opt}^B - 2 \cdot |\lambda_{d+1}|.$$

where λ_i is the i th largest eigenvalue of \mathbf{A} .

Using an appropriate pre-processing step and then running Algorithm 1 as a subroutine on a sub-sampled and low-rank version of \mathbf{A} , we output a k -subgraph \mathcal{Z}_d that has density opt_d . By essentially combining Lemmata 1 and 2 we obtain the following bounds.

Theorem 1. *The k -subgraph indexed by \mathcal{Z}_d has density*

$$\text{opt}_d \geq \text{opt}/(2 + o_n(1)) - 2 \cdot |\lambda_{d+1}|.$$

where λ_i is the i th largest, in magnitude, eigenvalue of the adjacency matrix of the graph. Moreover, if the graph is bipartite, or if the largest d eigenvalues of the graph are positive, then $\text{opt}_d \geq \text{opt} - 2 \cdot |\lambda_{d+1}|$.

Using bounds on eigenvalues of graphs, Theorem 1 translates to the following approximation guarantees.

Theorem 2. *If the densest- k -subgraph contains a constant fraction of all the edges, and $k = \Theta(\sqrt{E})$, then we can approximate DkS within a factor of $2 + \epsilon$, in time $n^{O(1/\epsilon^2)}$. If additionally the graph is bipartite, then we can approximate DkS within a factor of $1 + \epsilon$.*

Remark 2. *The above results are similar to the $1 + \epsilon$ ratio of (Arora et al., 1995), which holds for graphs where the densest- k -subgraph contains $\Omega(n^2)$ edges.*

Graph dependent bounds. For any given graph, after running our algorithm, we can compute an upper bound to the optimal density opt via bounds on opt^B , since it is easy to see that $\text{opt}^B \geq \text{opt}$. Our graph-dependent bound is the minimum of three upper bounds on the unknown optimal density:

Lemma 3. *The optimal density of DkS can be bounded as*

$$\text{opt} \leq \min \left\{ (1/k) \cdot \mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d} + |\lambda_{d+1}|, k - 1, \lambda_1 \right\}.$$

In our experimental section, we plot the above upper bounds, and show that for most tested graphs our algorithm performs *provably* within 70% from the upper bound on the optimal density. These are far stronger guarantees than the best available *a priori* bounds.

3. The Spannogram Framework

In this section, we describe how our constant rank solver operates by examining candidate vectors in a low-dimensional span of \mathbf{A} .

Here, we work on a rank- d matrix $\mathbf{A}_d = \mathbf{v}_1 \mathbf{u}_1^T + \dots + \mathbf{v}_d \mathbf{u}_d^T$ where $\mathbf{u}_i = \lambda_i \mathbf{v}_i$, and we wish to solve:

$$\max_{|\mathcal{X}|=|\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T (\mathbf{v}_1 \mathbf{u}_1^T + \dots + \mathbf{v}_d \mathbf{u}_d^T) \mathbf{1}_{\mathcal{Y}}. \quad (1)$$

Observe that we can rewrite (1) in the following way

$$\begin{aligned} & \max_{|\mathcal{X}|=|\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \left[\mathbf{v}_1 \cdot \underbrace{(\mathbf{u}_1^T \mathbf{1}_{\mathcal{Y}})}_{c_1} + \dots + \mathbf{v}_d \cdot \underbrace{(\mathbf{u}_d^T \mathbf{1}_{\mathcal{Y}})}_{c_d} \right] \\ &= \max_{|\mathcal{Y}|=k} \left(\max_{|\mathcal{X}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{v}_{\mathcal{Y}} \right), \end{aligned} \quad (2)$$

where $\mathbf{v}_{\mathcal{Y}} = \mathbf{v}_1 \cdot c_1 + \dots + \mathbf{v}_d \cdot c_d$ is an n -dimensional vector generated by the d -dimensional subspace spanned by $\mathbf{v}_1, \dots, \mathbf{v}_d$.

We will now make a key observation: for every fixed vector $\mathbf{v}_{\mathcal{Y}}$ in (2), the index set \mathcal{X} that maximizes $\mathbf{1}_{\mathcal{X}}^T \mathbf{v}_{\mathcal{Y}}$ can be easily computed. It is not hard to see that for any fixed vector $\mathbf{v}_{\mathcal{Y}}$, the k -subset \mathcal{X} that maximizes

$$\mathbf{1}_{\mathcal{X}}^T \mathbf{v}_{\mathcal{Y}} = \sum_{i \in \mathcal{X}} [\mathbf{v}_{\mathcal{Y}}]_i$$

corresponds to the set of k largest signed coordinates of $\mathbf{v}_{\mathcal{Y}}$. That is, the locally optimal k -set is $\text{top}_k(\mathbf{v}_{\mathcal{Y}})$.

We now wish to find all possible locally optimal sets \mathcal{X} . If we could possibly check all vectors $\mathbf{v}_{\mathcal{Y}}$, then we could find all locally optimal index sets $\text{top}_k(\mathbf{v}_{\mathcal{Y}})$.

Let us denote as \mathcal{S}_d the set of all k -subsets \mathcal{X} that are the optimal solutions of the inner maximization of (2) for any vector \mathbf{v} in the span of $\mathbf{v}_1, \dots, \mathbf{v}_d$

$$\mathcal{S}_d = \{\text{top}_k([\mathbf{v}_1 \cdot c_1 + \dots + \mathbf{v}_d \cdot c_d]) : c_1, \dots, c_d \in \mathbb{R}\}.$$

Clearly, this set contains all possible locally optimal \mathcal{X} sets of the form $\text{top}_k(\mathbf{v}_\mathcal{Y})$. Therefore, we can rewrite DBkS on \mathbf{A}_d as

$$\max_{|\mathcal{Y}|=k} \max_{\mathcal{X} \in \mathcal{S}_d} \mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}}. \quad (3)$$

The above problem can now be solved in the following way: for every set $\mathcal{X} \in \mathcal{S}_d$ find the locally optimal set \mathcal{Y} that maximizes $\mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}}$, that is, this will be $\text{top}_k(\mathbf{A}_d \mathbf{1}_{\mathcal{X}})$. Then, we simply need to test all such \mathcal{X}, \mathcal{Y} pairs on \mathbf{A}_d and keep the optimizer.

Due to the above, the problem of solving DBkS on \mathbf{A}_d is equivalent to constructing the set of k -supports \mathcal{S}_d , and then finding the optimal solution in that set. How large can \mathcal{S}_d be and can we construct it in polynomial time? Initially one could expect that the set \mathcal{S}_d could have size as big as $\binom{n}{k}$. Instead, we show that the set \mathcal{S}_d will be tremendously smaller.

Lemma 4. *The set \mathcal{S}_d has size at most $O(n^d)$ and can be built in time $O(n^{d+1})$ using Algorithm 2.*

3.1. Constructing the set \mathcal{S}_d

We build up to the general rank- d algorithm by explaining special cases that are easier to understand.

Rank-1 case. We start with the $d = 1$ case, where we have $\mathcal{S}_1 = \{\text{top}_k(c_1 \cdot \mathbf{v}_1) : c_1 \in \mathbb{R}\}$. It is not hard to see that there are only two supports to include in \mathcal{S}_1 : $\text{top}_k(\mathbf{v}_1)$ and $\text{top}_k(-\mathbf{v}_1)$. These two sets can be constructed in time in time $O(n)$, via a partial sorting and selection algorithm (Cormen et al., 2001). Hence, \mathcal{S}_1 has size 2 and can be constructed in time $O(n)$.

Rank-2 case. This is the first non-trivial d which exhibits the details of the Spannogram algorithm.

Let an auxiliary angle $\phi \in \Phi = [0, \pi)$ and let

$$\mathbf{c} = \begin{bmatrix} c_1 \\ c_2 \end{bmatrix} = \begin{bmatrix} \sin \phi \\ \cos \phi \end{bmatrix}. \quad ^3$$

Then, we re-express $c_1 \cdot \mathbf{v}_1 + c_2 \cdot \mathbf{v}_2$ in terms of ϕ as

$$\mathbf{v}(\phi) = \sin \phi \cdot \mathbf{v}_1 + \cos \phi \cdot \mathbf{v}_2. \quad (4)$$

This means that we can rewrite the set \mathcal{S}_2 as:

$$\mathcal{S}_2 = \{\text{top}_k(\pm \mathbf{v}(\phi)), \phi \in [0, \pi)\}.$$

Observe that each element of $\mathbf{v}(\phi)$ is a continuous *spectral curve* in ϕ : $[\mathbf{v}(\phi)]_i = [\mathbf{v}_1]_i \sin(\phi) + [\mathbf{v}_2]_i \cos(\phi)$. Consequently, the top/bottom- k supports of $\mathbf{v}(\phi)$ (*i.e.*, $\text{top}_k(\pm \mathbf{v}(\phi))$) are themselves a function of ϕ . How can we find all possible supports?

³Observe that when we scan ϕ , the vectors $\mathbf{c}, -\mathbf{c}$ express all possible unit norm vectors on the circle.

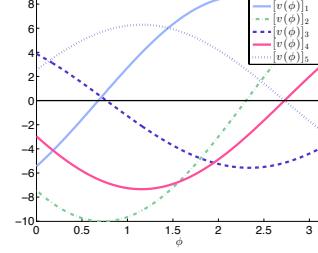


Figure 1. A rank $d = 2$ spannogram for $n = 5$ and two random vectors $\mathbf{v}_1, \mathbf{v}_2$. Observe that every two curves intersect in exactly one point. These intersection points define intervals in which a top- k set is invariant.

The Spannogram. In Fig. 1, we draw an example plot of five curves $[\mathbf{v}(\phi)]_i$, $i = 1, \dots, 5$, which we call a spannogram. From the spannogram in Fig. 1, we can see that the continuity of these sinusoidal curves implies a “local invariance” property of the top/bottom k supports $\text{top}_k(\pm \mathbf{v}(\phi))$, in a small neighborhood around a fixed ϕ . So, when does a top/bottom- k support change? The index sets $\text{top}_k(\pm \mathbf{v}(\phi))$ change *if and only if* two curves cross, *i.e.*, when the ordering of two elements $[\mathbf{v}(\phi)]_i, [\mathbf{v}(\phi)]_j$ changes.

Finding all supports: There are n curves and each pair intersects at exactly one point in the Φ domain⁴. Therefore, there are exactly $\binom{n}{2}$ intersection points. These $\binom{n}{2}$ intersection points define $\binom{n}{2} + 1$ intervals. Within an interval the top/bottom k supports $\text{top}_k(\pm \mathbf{v}(\phi))$ remain the same. Hence, it is now clear that $|\mathcal{S}_2| \leq 2 \binom{n}{2} = O(n^2)$.

A way to find all supports in \mathcal{S}_2 is to compute the $\mathbf{v}(\phi_{i,j})$ vectors on the intersection points of two curves i, j , and then the supports in the two adjacent intervals of such intersection point. The $\mathbf{v}(\phi_{i,j})$ vector on an intersection point of two curves i and j can be easily computed by first solving a set of linear equations $[\mathbf{v}(\phi_{i,j})]_i = [\mathbf{v}(\phi_{i,j})]_j \Rightarrow (\mathbf{e}_i - \mathbf{e}_j)^T [\mathbf{v}_1 \ \mathbf{v}_2] \mathbf{c}_{i,j} = \mathbf{0}_{2 \times 1}$ for the unknown vector $\mathbf{c}_{i,j}$, where \mathbf{e}_i is the i -th column of the $n \times n$ identity matrix, *i.e.*,

$$\mathbf{c}_{i,j} = \text{nullspace}((\mathbf{e}_i - \mathbf{e}_j)^T [\mathbf{v}_1 \ \mathbf{v}_2]).$$

Then, we compute $\mathbf{v}(\phi_{i,j}) = [\mathbf{v}_1 \ \mathbf{v}_2] \mathbf{c}_{i,j}$. Further details on breaking ties in $\text{top}_k(\mathbf{v}(\phi_{i,j}))$ can be found in the supplemental material.

Computational cost: We have $\binom{n}{2}$ intersection points, where we calculate the top/bottom k supports for each $\mathbf{v}(\phi_{i,j})$. The top/bottom k elements of every $\mathbf{v}(\phi_{i,j})$ can be computed in time $O(n)$ using a partial sorting

⁴Here we assume that the curves are in *general position*. This can be always accomplished by infinitesimally perturbing the curves as in (Papailiopoulos et al., 2013).

and selection algorithm (Cormen et al., 2001). Since we perform this routine a total of $O(\binom{n}{2})$ times, the total complexity of our rank-2 algorithm is $O(n^3)$.

General Rank- d case. The algorithm generalizes to arbitrary dimension d , as we show in the supplemental material; its pseudo-code is given as Algorithm 2.

Remark 3. *Observe that the computation of each loop in line 2 of Algorithm 2 can be computed in parallel. This will allow us to parallelize the Spannogram.*

Algorithm 2 Spannogram(k, \mathbf{V}_d)

```

1:  $\mathcal{S}_d = \emptyset$ 
2: for all  $(i_1, \dots, i_d) \in \{1, \dots, n\}^d$  and  $s \in \{-1, 1\}$  do
3:    $\mathbf{c} = s \cdot \text{nullspace} \begin{pmatrix} [(\mathbf{V}_d)_{i_1, \dots, -[\mathbf{V}_d]_{i_2, :}} \\ \vdots \\ [\mathbf{V}_d]_{i_1, :} - [\mathbf{V}_d]_{i_d, :} ] \end{pmatrix}$ 
4:    $\mathbf{v} = \mathbf{V}_d^T \mathbf{c}$ 
5:    $\mathcal{S} = \text{top}_k(\mathbf{v})$ 
6:    $\mathcal{T} = \mathcal{S} - \{i_1, \dots, i_d\}$ 
7:   for all  $\binom{d}{d-|\mathcal{T}|}$  subsets  $\mathcal{J}$  of  $(i_1, \dots, i_d)$  do
8:      $\mathcal{S}_d = \mathcal{S}_d \cup (\mathcal{T} \cup \mathcal{J})$ 
9:   end for
10: end for
11: Output:  $\mathcal{S}_d$ .

```

3.2. An approximate \mathcal{S}_d in nearly-linear time

In our exact solver, we solve DBkS on \mathbf{A}_d in time $O(n^{d+1})$. Surprisingly, when \mathbf{A}_d has only positive eigenvalues, then we can tightly approximate DBkS on \mathbf{A}_d in nearly linear time.

Theorem 3. *Let the d largest eigenvalues of the graph be positive and let $\lambda_d/\lambda_{d+1} = \Theta(1)$. Then, using Algorithm 3 as a subroutine in Algorithm 1 yields, with high probability, a subgraph with density $(1 - \frac{1}{\log n}) \cdot \text{opt}_d$ in nearly-linear time $O(m \cdot \log n + n \cdot \log^{d+1} n)$, where m is the number of edges.*

The main idea is that instead of checking all $O(n^d)$ possible k sets in \mathcal{S}_d , we can approximately solve the problem by randomly sampling $(\log^d n)$ vectors in the span of $\mathbf{v}_1, \dots, \mathbf{v}_d$. Our proof is based on the fact that we can “approximate” the surface of the d -dimensional sphere with $M = O(\log^{d+1} n)$ randomly sampled vectors from the span of $\mathbf{v}_1, \dots, \mathbf{v}_d$, which then allows us to identify, with high probability, near-optimal candidates in \mathcal{S}_d . The modified algorithm is very simple and is given below; its analysis can be found in the supplemental material.

Algorithm 3 Spannogram_approx($k, \mathbf{V}_d, \Lambda_d$)

```

1: for  $i = 1 : d \cdot 2^d \cdot \log^{d+1} n$  do
2:    $\mathbf{v} = (\Lambda_d^{1/2} \cdot \mathbf{V}_d)^T \cdot \text{randn}(d, 1)$ 
3:    $\mathcal{S}_d = \mathcal{S}_d \cup \text{top}_k(\mathbf{v}) \cup \text{top}_k(-\mathbf{v})$ 
4: end for
5: Output:  $\mathcal{S}_d$ .

```

4. Scaling up

In this section, we present the two key scalability features that allow us to scale up to graphs with billions of edges.

4.1. Vertex Sparsification

We introduce a very simple and efficient pre-processing step for discarding vertices that are unlikely to appear in a top k set in \mathcal{S}_d . This step runs after we compute \mathbf{A}_d and uses the leverage score, $\ell_i = \sum_{j=1}^d [\mathbf{V}_d]_{i,j}^2 |\lambda_j|$, of the i -th vertex to decide whether we will discard it or not. We show in the supplemental material, that by appropriately setting a threshold, we can guarantee a provable bound on the error introduced. In our experimental results, the above elimination is able to reduce n to approximately $\hat{n} \approx 10 \cdot k$ for a provably small additive error, even for data sets where $n = 10^8$.

4.2. MapReduce Implementation

A MapReduce implementation allows scaling out to a large number of compute nodes that can work in parallel. The reader can refer to (Meng & Mahoney; Bahmani et al., 2012)) for a comprehensive treatment of the MapReduce paradigm. In short, the Hadoop/MapReduce infrastructure stores the input graph as a distributed file spread across multiple machines; it provides a tuple streaming abstraction, where each `map` and `reduce` function receives and emits tuples as $(key, value)$ pairs. The role of the keys is to ensure information aggregation: all the tuples with the same key are processed by the same reducer.

For the spectral decomposition step of our scheme we design a simple implementation of the power method in MapReduce. The details are beyond the scope of this work; high-performance implementations are already available in the literature, e.g. (Lin & Schatz, 2010). We instead focus on the novel implementation of the Spannogram.

Our MapReduce implementation of the rank-2 Spannogram is outlined in Algorithm 4. The Mapper is responsible for the duplication and dissemination of the eigenvectors, $\mathbf{V}_2, \mathbf{U}_2 = \mathbf{V}_2 \Lambda_2$, to all reducers. Line 3 emits the j -th row of \mathbf{V}_2 and \mathbf{U}_2 once for every node i . Since i is used as the key, this ensures that every reducer receives $\mathbf{V}_2, \mathbf{U}_2$ in their entirety.

From the breakdown of the Spannogram in Section 3, it is understood that, for the rank-2 case, it suffices to solve a simple system of equations for every pair of nodes. The Reducer for node i receives the full eigenvectors $\mathbf{V}_2, \mathbf{U}_2$ and is responsible for solving the problem for every pair (i, j) , where $j > i$. Then, Line 6 emits the best candidate computed at Reducer i . A

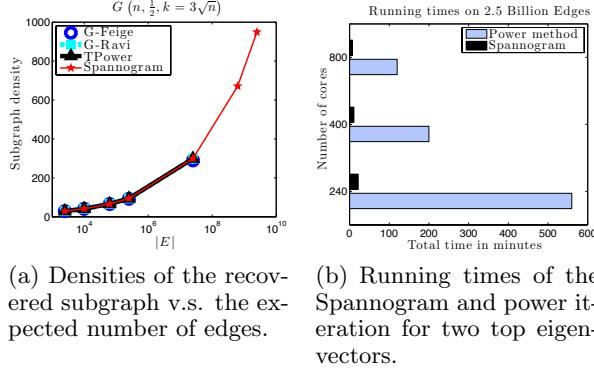


Figure 2. Planted clique experiments for random graphs.

trivial final step, not outlined here, collects all n^2 candidate sets and keeps the best one as the final solution.

The basic outline in Algorithm 4 comes with heavy communication needs and was chosen here for ease of exposition. The more efficient version that we implement, does not replicate $\mathbf{V}_2, \mathbf{U}_2$ n times. Instead, the number of reducers – say $R = n^\alpha$ – is fine-tuned to the capabilities of the cluster. The mappers emit $\mathbf{V}_2, \mathbf{U}_2 R$ times, once for every reducer. Then, reducer r is responsible for solving for node pairs (i, j) , where $i \equiv r \pmod{R}$ and $j > i$. Depending on the performance bottleneck, different choices for α are more appropriate. We divide the construction of the $O(n^2)$ candidate sets in \mathcal{S}_2 to $O(n^\alpha)$ reducers and each of them computes $O(n^{2-\alpha})$ candidate subgraphs. The total communication cost for this parallelization scheme is $O(n^{1+\alpha})$: n^α reducers need to have access to the entire $\mathbf{V}_2, \mathbf{U}_2$ that has $2 \cdot 2 \cdot n$ entries. Moreover, the total computation cost for each reducer is $O(n^{3-\alpha})$.

Algorithm 4 Spannogram_{MR}($\mathbf{V}_2, \mathbf{U}_2$)

```

1: Map( $\{[\mathbf{V}_2]_{j,:}, [\mathbf{U}_2]_{j,:}, j\}$ ):
2: for  $i = 1 : n$  do
3:   emit:  $\langle i, \{[\mathbf{V}_2]_{j,1}, [\mathbf{V}_2]_{j,2}, [\mathbf{U}_2]_{j,1}, [\mathbf{U}_2]_{j,2}, j\} \rangle$ 
4: end for
1: Reduce $i$ ( $\langle i, \{[\mathbf{V}_2]_{j,1}, [\mathbf{V}_2]_{j,2}, [\mathbf{U}_2]_{j,1}, [\mathbf{U}_2]_{j,2}, j\} \rangle, \forall j \rangle$ :
2: for each  $j \geq i + 1$  do
3:    $\mathbf{c} = \text{nullspace}([\mathbf{V}]_{i,:} - [\mathbf{V}]_{j,:})$ 
4:    $[\mathbf{den}_j, \{\mathcal{X}_j, \mathcal{Y}_j\}] = \max_{|\mathcal{Y}|=k, \mathcal{X} \in \text{top}_k(\pm \mathbf{V}_2 \mathbf{c})} \mathbf{1}_{\mathcal{X}} \mathbf{V}_2 \mathbf{U}_2^T \mathbf{1}_{\mathcal{Y}}$ 
5: end for
6: emit:  $\langle i, \{\mathcal{X}_i, \mathcal{Y}_i\} = \max_j \mathbf{1}_{\mathcal{X}_j} \mathbf{V}_2 \mathbf{U}_2^T \mathbf{1}_{\mathcal{Y}_j} \rangle$ 

```

5. Experimental Evaluation

We experimentally evaluate the performance of our algorithm and compare it to the truncated power method (TPower) of (Yuan & Zhang, 2011), a greedy algorithm by (Feige et al., 2001) (GFeige) and another greedy algorithm by (Ravi et al., 1994) (GRavi). We performed experiments on synthetic dense subgraphs

and also massive real graphs from multiple sources. In all experiments we compare the density of the subgraph obtained by the Spannogram to the density of the output subgraphs given by the other algorithms.

Our experiments illustrate three key points: (1) for all tested graphs, our method outperforms – some times significantly – all other algorithms compared; (2) our method is highly scalable, allowing us to solve far larger problem instances; (3) our data-dependent upper bound in many cases provide a certificate of near-optimality, far more accurate and useful, than what *a priori* bounds are able to do.

Planted clique. We first consider the so-called (and now much studied) Planted Clique problem: we seek to find a clique of size k that has been planted in a graph where all other edges are drawn independently with probability $1/2$. We scale our randomized experiments from $n = 100$ up to 10^5 . In all cases we set the size of the clique to $k = 3 \cdot \sqrt{n}$ – close to what is believed to be the critical computability threshold. In all our experiments, GRavi, TPower, and the Spannogram successfully recovered the hidden clique. However, as can be seen in Fig. 2, the Spannogram algorithm is the only one able to scale up to $n = 10^5$ – a massive dense graph with about 2.5 billion edges. The reason is that this graph does not fit in the main memory of one machine and caused all centralized algorithms to crash after several hours. Our MapReduce implementation scales out smoothly, since it splits the problem over multiple smaller problems solved in parallel.

Specifically, we used Amazon Wireless Services’ Elastic MapReduce framework ([aws](#)). We implemented our `map` and `reduce` functions in Python and used the MR-Job class ([mrj](#)). For our biggest experiments we used a 100-machine strong cluster, consisting of m1.large AWS instances (a total of 800 cores).

The running times of our experiments over MapReduce are shown in Fig. 2(b). The main bottleneck is the computation of the first two eigenvectors which is performed by repeating the power iteration for few (typically 4) iterations. This step is not the emphasis of this work and has not been optimized. The Spannogram algorithm is significantly faster and the benefits of parallelization are clear since it is CPU intensive.

In principle, the other algorithms could be also implemented over MapReduce, but that requires non-trivial distributed algorithm design. As is well-known, *e.g.*, ([Meng & Mahoney](#)), implementing iterative machine learning algorithms over MapReduce can be a significant task and schemes which perform worse in standard metrics can be highly preferable for this parallel

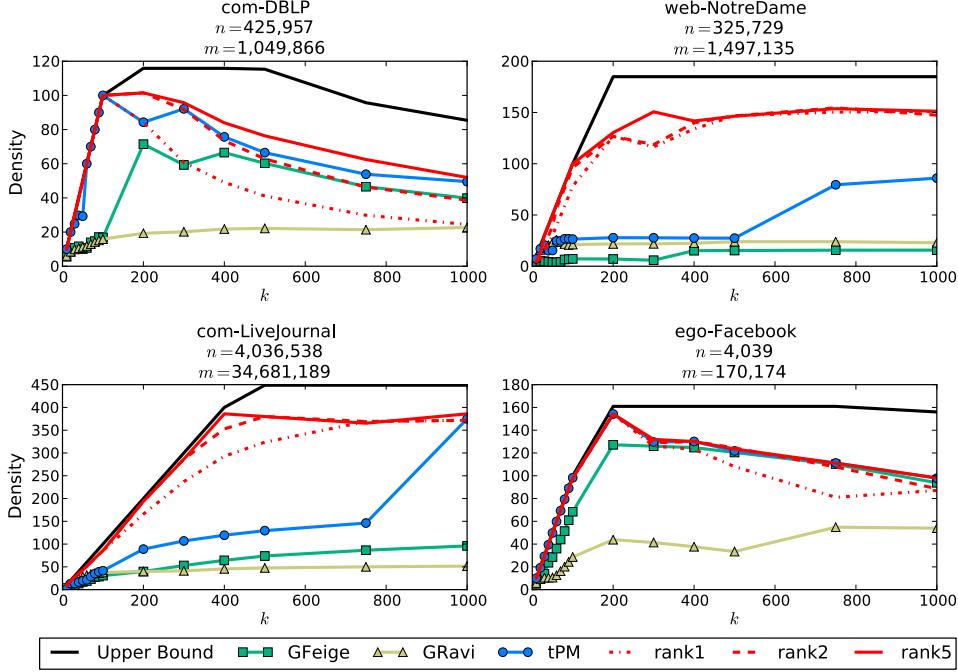


Figure 3. Subgraph density vs. subgraph size (k). We compare our DkS Spannogram algorithm with the algorithms from (Feige et al., 2001) (GFeige), (Ravi et al., 1994) (GRavi), and (Yuan & Zhang, 2011) (tPM). Across all subgraph sizes k , we obtain higher subgraph densities using Spannograms of rank $d = 2$ or 5. We also obtain a *provable data-dependent upper bound* (solid black line) on the objective. This proves that for these data sets, our algorithm is typically within 80% from optimality, for all sizes up to $k = 250$, and indeed for small subgraph sizes we find a clique which is clearly optimal. Further experiments on multiple other data sets are shown in the supplemental material.

framework. Careful MapReduce algorithmic design is needed especially for dense graphs like the one in the hidden clique problem.

Real Datasets. Next, we demonstrate our method’s performance in real datasets and also illustrate the power of our data-dependent bounds. We run experiments on large graphs from different applications and our findings are presented in Fig. 3. The figure compares the density achieved by the Spannogram algorithm for rank 1, 2 and 5 to the performance of GFeige, GRavi and TPower. The figure shows that the rank-2 and rank-5 versions of our algorithm, improve – sometimes significantly – over the other techniques. Our novel data-dependent upper-bound shows that our results on these data sets are provably near-optimal.

The experiments are performed for two community graphs (com-LiveJournal and com-DBLP), a web graph (web-NotreDame), and a subset of the Facebook graph. A larger set of experiments is included in the supplemental material. Note that the largest graph in Figure 3 contains no more than 35 million edges; these cases fit in the main memory of a single machine and

the running times are presented in the supplemental material, all performed on a standard Macbook Pro laptop using Matlab. In summary, rank-2 took less than one second for all these graphs while prior work methods took approximately the same time, up to a few seconds. Rank-1 was significantly faster than all other methods in all tested graphs and took fractions of a second. Rank-5 took up to 1000 seconds for the largest graph (LiveJournal).

We conclude that our algorithm is an efficient option for finding dense subgraphs. Different rank choices give a tradeoff between accuracy and performance while the parallel nature allows scalability when needed. Further, our theoretical upper-bound can be useful for practitioners investigating dense structures in large graphs.

6. Acknowledgments

The authors would like to acknowledge support from NSF grants CCF 1344364, CCF 1344179, DARPA XDATA, and research gifts by Google, Docomo and Microsoft.

References

- Amazon Web Services, Elastic Map Reduce. URL <http://aws.amazon.com/elasticmapreduce/>.
- MRJob. URL <http://pythonhosted.org/mrjob/>.
- Alon, Noga, Lee, Troy, Shraibman, Adi, and Vempala, Santosh. The approximate rank of a matrix and its algorithmic applications: approximate rank. In *Proceedings of the 45th annual ACM symposium on Symposium on theory of computing*, pp. 675–684. ACM, 2013.
- Ames, Brendan PW. *Convex relaxation for the planted clique, biclique, and clustering problems*. PhD thesis, University of Waterloo, 2011.
- Arora, Sanjeev, Karger, David, and Karpinski, Marek. Polynomial time approximation schemes for dense instances of np-hard problems. In *STOC*, 1995.
- Asahiro, Yuichi, Iwama, Kazuo, Tamaki, Hisao, and Tokuyama, Takeshi. Greedily finding a dense subgraph. *Journal of Algorithms*, 34(2):203–221, 2000.
- Asteris, Megasthenis, Papailiopoulos, Dimitris S, and Karystinos, George N. The sparse principal component of a constant-rank matrix. *IEEE Trans. IT*, 60(4):228–2290, 2014.
- Bahmani, Bahman, Kumar, Ravi, and Vassilvitskii, Sergei. Densest subgraph in streaming and mapreduce. *Proceedings of the VLDB Endowment*, 5(5):454–465, 2012.
- Bhaskara, Aditya, Charikar, Moses, Chlamtac, Eden, Feige, Uriel, and Vijayaraghavan, Aravindan. Detecting high log-densities: an $O(n^{1/4})$ approximation for densest k-subgraph. In *STOC*, 2010.
- Boutsidis, Christos, Mahoney, Michael W, and Drineas, Petros. An improved approximation algorithm for the column subset selection problem. In *Proceedings of the twentieth Annual ACM-SIAM Symposium on Discrete Algorithms*, pp. 968–977. Society for Industrial and Applied Mathematics, 2009.
- Cormen, Thomas H, Leiserson, Charles E, Rivest, Ronald L, and Stein, Clifford. *Introduction to algorithms*. MIT press, 2001.
- d’Aspremont, Alexandre et al. Weak recovery conditions using graph partitioning bounds. 2010.
- Dourisboure, Yon, Geraci, Filippo, and Pellegrini, Marco. Extraction and classification of dense communities in the web. In *WWW*, 2007.
- Feige, Uriel and Langberg, Michael. Approximation algorithms for maximization problems arising in graph partitioning. *Journal of Algorithms*, 41(2):174–211, 2001.
- Feige, Uriel, Peleg, David, and Kortsarz, Guy. The dense k-subgraph problem. *Algorithmica*, 29(3):410–421, 2001.
- Gibson, David, Kumar, Ravi, and Tomkins, Andrew. Discovering large dense subgraphs in massive graphs. In *PVLDB*, 2005.
- Hu, Haiyan, Yan, Xifeng, Huang, Yu, Han, Jiawei, and Zhou, Xianghong Jasmine. Mining coherent dense subgraphs across massive biological networks for functional discovery. *Bioinformatics*, 21(suppl 1):i213–i221, 2005.
- Jethava, Vinay, Martinsson, Anders, Bhattacharyya, Chiranjib, and Dubhashi, Devdatt. The lovasz theta function, svms and finding large dense subgraphs. In *NIPS*, 2012.
- Karystinos, George N and Liavas, Athanasios P. Efficient computation of the binary vector that maximizes a rank-deficient quadratic form. *IEEE Trans. IT*, 56(7):3581–3593, 2010.
- Khot, Subhash. Ruling out ptas for graph min-bisection, densest subgraph and bipartite clique. In *FOCS*, 2004.
- Lin, Jimmy and Schatz, Michael. Design patterns for efficient graph algorithms in mapreduce. In *Proceedings of the Eighth Workshop on Mining and Learning with Graphs*, pp. 78–85. ACM, 2010.
- Mahoney, Michael W and Drineas, Petros. Cur matrix decompositions for improved data analysis. *Proceedings of the National Academy of Sciences*, 106(3):697–702, 2009.
- Meng, Xiangrui and Mahoney, Michael W. Robust regression on mapreduce. *ICML 2013*, (to appear).
- Miller, B, Bliss, N, and Wolfe, P. Subgraph detection using eigenvector l1 norms. In *NIPS*, 2010.
- Papailiopoulos, Dimitris S, Dimakis, Alexandros G, and Korokythakis, Stavros. Sparse pca through low-rank approximations. *arXiv preprint arXiv:1303.0551*, 2013.
- Ravi, Sekharipuram S, Rosenkrantz, Daniel J, and Tayi, Giri K. Heuristic and special case algorithms for dispersion problems. *Operations Research*, 42(2):299–310, 1994.
- Saha, Barna, Hoch, Allison, Khuller, Samir, Raschid, Louiza, and Zhang, Xiao-Ning. Dense subgraphs with restrictions and applications to gene annotation graphs. In *Research in Computational Molecular Biology*, pp. 456–472. Springer, 2010.
- Srivastav, Anand and Wolf, Katja. *Finding dense subgraphs with semidefinite programming*. Springer, 1998.
- Suzuki, Akiko and Tokuyama, Takeshi. Dense subgraph problems with output-density conditions. In *Algorithms and Computation*, pp. 266–276. Springer, 2005.
- Yuan, Xiao-Tong and Zhang, Tong. Truncated power method for sparse eigenvalue problems. *arXiv preprint arXiv:1112.2679*, 2011.

Supplemental Material for:
Finding Dense Subgraphs via Low-Rank Bilinear Optimization

055
056
057
058
059
060
061
062
063
064
065
066
067
068
069
070
071
072
073
074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
102
103
104
105
106
107
108
109

1. Proof of Lemma 4: Building the set \mathcal{S}_d for arbitrary d -dimensional subspaces

In our general case, we solve DBkS on

$$\mathbf{A}_d = \mathbf{V}_d \mathbf{U}_d^T = \sum_{i=1}^d \mathbf{v}_i \mathbf{u}_i^T$$

where

$$\mathbf{V}_d = [v_1 \dots v_d] \text{ and } \mathbf{U}_d = [\lambda_1 \cdot v_1 \dots \lambda_d \cdot v_d].$$

Solving the problem on \mathbf{A}_d is equivalent to answering the following combinatorial question:

"how many different top- k supports are there in a d -dimensional subspace: $\text{top}_k(c_1 \cdot \mathbf{v}_1 + \dots + c_d \cdot \mathbf{v}_d)$?"

Here we define $d - 1$ auxiliary angles $\phi_1, \dots, \phi_{d-1} \in \Phi = [0, \pi]$ and we rewrite the coefficients c_1, \dots, c_d as

$$\mathbf{c} = \begin{bmatrix} c_1 \\ \vdots \\ c_d \end{bmatrix} = \begin{bmatrix} \sin \phi_1 \\ \cos \phi_1 \sin \phi_2 \\ \vdots \\ \cos \phi_1 \cos \phi_2 \dots \sin \phi_{d-1} \\ \cos \phi_1 \cos \phi_2 \dots \cos \phi_{d-1} \end{bmatrix}.$$

Clearly we can express every vector in the span of \mathbf{V}_d as a linear combination $c_1 \cdot \mathbf{v}_1 + \dots + c_d \cdot \mathbf{v}_d$ in terms of ϕ :

$$\mathbf{v}(\phi_1, \dots, \phi_{d-1}) = (\sin \phi_1) \cdot \mathbf{v}_1 + (\cos \phi_1 \sin \phi_2) \cdot \mathbf{v}_2 + \dots + (\cos \phi_1 \cos \phi_2 \dots \cos \phi_{d-1}) \cdot \mathbf{v}_d. \quad (1)$$

For notation simplicity let us define a vector that contains all $d - 1$ auxiliary phase variables

$$\varphi = [\phi_1, \dots, \phi_{d-1}].$$

We can use the above derivations to rewrite the set \mathcal{S}_d that contains all top k coordinates in the span of \mathbf{V}_d as:

$$\begin{aligned} \mathcal{S}_d &= \{\text{top}_k(c_1 \cdot \mathbf{v}_1 + \dots + c_d \cdot \mathbf{v}_d) : c_1, \dots, c_d \in \mathbb{R}\} \\ &= \{\text{top}_k \pm (\mathbf{v}(\varphi)) : \varphi \in \Phi^{d-1}\} \\ &= \{\text{top}_k \pm ((\sin \phi_1) \cdot \mathbf{v}_1 + (\cos \phi_1 \sin \phi_2) \cdot \mathbf{v}_2 + \dots + (\cos \phi_1 \cos \phi_2 \dots \cos \phi_{d-1}) \cdot \mathbf{v}_d), \varphi \in \Phi^{d-1}\} \end{aligned}$$

Observe again that each element of $\mathbf{v}(\varphi)$ is a continuous *spectral curve* in the $d - 1$ auxiliary variables:

$$[\mathbf{v}(\varphi)]_i = (\sin \phi_1) \cdot [\mathbf{v}_1]_i + (\cos \phi_1 \sin \phi_2) \cdot [\mathbf{v}_2]_i + \dots + (\cos \phi_1 \cos \phi_2 \dots \cos \phi_{d-1}) \cdot [\mathbf{v}_d]_i.$$

Consequently, the top/bottom- k supports of $\mathbf{v}(\varphi)$ (i.e., $\text{top}_k(\pm \mathbf{v}(\varphi))$) are themselves a function of the $d - 1$ variables in φ . How can we find all possible supports?

Remark 1. In our general problem we wish to find all top and bottom k coordinates that appear in a d -dimensional subspace. In the following discussion, for simplicity we handle the top k coordinates problem. Finding the bottom k trivially follows, by just checking the smallest k coordinates of each vector $c_1 \cdot \mathbf{v}_1 + \dots + c_d \cdot \mathbf{v}_d$ that we construct using our algorithm.

110 1.1. Ranking regions for a single coordinate $[\mathbf{v}(\varphi)]_i$

165

 111 We now show that for each single coordinate $[\mathbf{v}(\varphi)]_i$, we can partition Φ^{d-1} in regions, wherein the i th coordinate
 112 $[\mathbf{v}(\varphi)]_i$ retains the same ranking relative to the other $n - 1$ coordinates in the vector $\mathbf{v}(\varphi)$.

166

167

 114 Let us first consider for simplicity $[\mathbf{v}(\varphi)]_1$. We aim to find all values of φ where $[\mathbf{v}(\varphi)]_1$ is in one of the the k
 115 largest coordinates of $\mathbf{v}(\varphi)$. We observe that this region can be characterized by using n boundary tests:

169

170

$$\begin{aligned} [\mathbf{v}(\varphi)]_1 &\geq [\mathbf{v}(\varphi)]_2 \\ [\mathbf{v}(\varphi)]_1 &\geq [\mathbf{v}(\varphi)]_3 \\ &\vdots \\ [\mathbf{v}(\varphi)]_1 &\geq [\mathbf{v}(\varphi)]_n \end{aligned}$$

171

172

173

174

 123 Each of the above boundary tests defines a bounding curve that partitions the Φ^{d-1} domain. We refer to this
 124 bounding curve as $\mathcal{B}_{1,j}(\varphi) : \Phi^{d-1} \mapsto \Phi^{d-2}$. A $\mathcal{B}_{1,j}(\varphi)$ curve partitions Φ and defines two regions of φ angles:

175

176

177

$$\mathcal{R}_{1>j} = \{\varphi \in \Phi^{d-1} : [\mathbf{v}(\varphi)]_1 > [\mathbf{v}(\varphi)]_j\} \text{ and } \mathcal{R}_{1\leq j} = \{\varphi \in \Phi^{d-1} : [\mathbf{v}(\varphi)]_1 \leq [\mathbf{v}(\varphi)]_j\} \quad (2)$$

181

182

 128 such that $\mathcal{R}_{1>j} \cup \mathcal{R}_{1\leq j} = \Phi^{d-1}$.

183

 129 Observe that these $n - 1$ curves $\mathcal{B}_{1,1}(\varphi), \dots, \mathcal{B}_{1,n}(\varphi)$ partition Φ in disjoint cells, $\mathcal{C}_1^1, \dots, \mathcal{C}_T^1$, such that

184

185

$$\bigcup_{i=1}^T \mathcal{C}_i^1 = \Phi^{d-1}.$$

186

187

188

189

 135 Within each cell \mathcal{C}_i^1 , the first coordinate $[\mathbf{v}(\varphi)]_1$ retains a fixed ranking relative to the rest of the elements in
 136 $\mathbf{v}(\varphi)$, e.g., for a specific cell it might be the largest element, and in another cell it might be the 10th smallest, etc.
 137 This happens because for all values of φ in a single cell, the respective ordering $[\mathbf{v}(\varphi)]_1 \geq [\mathbf{v}(\varphi)]_2, \dots, [\mathbf{v}(\varphi)]_1 \geq$
 138 $[\mathbf{v}(\varphi)]_n$ remains the same.

190

191

192

193

 139 If we have access to a single point, say φ_0 , that belongs to a specific cell, say \mathcal{C}_j^1 , then we can calculate $[\mathbf{v}(\varphi_0)]_1$
 140 and find the ranking of the first coordinate $[\mathbf{v}(\varphi)]_1$, that remains invariant for all $\varphi \in \mathcal{C}_j^1$. Hence, if we visit all
 141 these cells, then we can find all possible rankings that the first coordinate $[\mathbf{v}(\varphi)]_1$ takes in the d -dimensional span
 142 of $\mathbf{v}_1, \dots, \mathbf{v}_d$. In the following subsections, we show that the number of these cells is bounded by $T \leq 2^d \binom{n-1}{d-1}$.

194

195

196

197

 144 Observe that each bounding curve $\mathcal{B}_{1,i}(\varphi)$ has a one-to-one correspondence to an equation $[\mathbf{v}(\varphi)]_1 = [\mathbf{v}(\varphi)]_j$,
 145 which is linear in \mathbf{c} :

198

199

200

201

$$[\mathbf{v}(\varphi)]_1 = [\mathbf{v}(\varphi)]_j \Rightarrow \mathbf{e}_1^T \mathbf{V}_d \mathbf{c} - \mathbf{e}_j^T \mathbf{V}_d \mathbf{c} = 0 \Rightarrow (\mathbf{e}_1 - \mathbf{e}_j)^T \mathbf{V}_d \mathbf{c} = 0. \quad (3)$$

202

203

 146 Due to their linear characterization with respect to \mathbf{c} , it is easy to see that each $(d - 1)$ -tuple of bounding curves
 147 intersects on a single point in Φ^{d-1} .¹

204

205

206

$$\begin{aligned} [\mathbf{v}(\varphi)]_1 &= [\mathbf{v}(\varphi)]_{i_1} & (\mathbf{e}_1 - \mathbf{e}_{i_1})^T \mathbf{V}_d \mathbf{c} = 0 \\ [\mathbf{v}(\varphi)]_1 &= [\mathbf{v}(\varphi)]_{i_2} & (\mathbf{e}_1 - \mathbf{e}_{i_2})^T \mathbf{V}_d \mathbf{c} = 0 \\ &\vdots & \vdots \\ [\mathbf{v}(\varphi)]_1 &= [\mathbf{v}(\varphi)]_{i_{d-1}} & (\mathbf{e}_1 - \mathbf{e}_{i_{d-1}})^T \mathbf{V}_d \mathbf{c} = 0 \end{aligned} \Rightarrow \begin{bmatrix} (\mathbf{e}_1 - \mathbf{e}_{i_1})^T \\ (\mathbf{e}_1 - \mathbf{e}_{i_2})^T \\ \vdots \\ (\mathbf{e}_1 - \mathbf{e}_{i_{d-1}})^T \end{bmatrix} \mathbf{V}_d \mathbf{c} = \mathbf{0}_{(d-1) \times 1}.$$

207

208

209

210

211

 151 Let us denote the solution of the above linear inverse problem as $\mathbf{c}_{1,i_1, \dots, i_{d-1}}$. We refer to $\mathbf{c}_{1,i_1, \dots, i_{d-1}}$ as an
 152 intersection vector. For each intersection vector $\mathbf{c}_{1,i_1, \dots, i_{d-1}}$, we can compute its polar expression and solve for
 153 the angles φ that generate it. These $d - 1$ input angles correspond exactly to the intersection point of $d - 1$ curves
 154 specified by the above $d - 1$ equations. We denote these $d - 1$ angles that generate $\mathbf{c}_{1,i_1, \dots, i_{d-1}}$, as $\varphi_{1,i_1, \dots, i_{d-1}}$
 155 which we refer to as the intersection point of the $d - 1$ curves $\mathcal{B}_{1,i_1}(\varphi), \dots, \mathcal{B}_{1,i_{d-1}}(\varphi)$.

212

213

214

215

216

 162 ¹as a matter of fact, due to the sign ambiguity of the solution, this corresponds to two intersection points. However,
 163 the following discussion omits this technical detail for simplicity.

217

218

219

220	Since, the $\varphi_{1,i_1,\dots,i_{d-1}}$ intersection points are defined for every $d-1$ curves, the total number of intersection points	275
221	is $\binom{n-1}{d-1}$. In the following subsections, we show how we can visit all cells by just examining these intersection	276
222	points.	277
223	We proceed to show that if we visit the adjacent cells of the intersection points defined for all coordinates, then	278
224	we can find all top- k supports in the span of \mathbf{V}_d .	279
225		280
226		281
227	1.2. Visiting all cells = finding all top k supports	282
228	Our goal is to find all top- k supports that can appear in the span of \mathbf{V}_d . To do so, it is sufficient to visit the	283
229	cells where $[\mathbf{v}(\varphi)]_1$ is the k -th largest coordinate, then the cells where $[\mathbf{v}(\varphi)]_2$ is the k -largest, and so on. Within	284
230	such cells, one coordinate (say $[\mathbf{v}(\varphi)]_i$) remains always the k -th largest, while the identities of the bottom $n-k$	285
231	coordinates remain the same. This means that in such a cell, we have that	286
232		287
233	$[\mathbf{v}(\varphi)]_i \geq [\mathbf{v}(\varphi)]_{j_1}, \dots, [\mathbf{v}(\varphi)]_i \geq [\mathbf{v}(\varphi)]_{j_{n-k}}$	288
234		289
235	for all φ in that cell and some specific $n-k$ other coordinates indexed by j_1, \dots, j_{n-k} . Hence, although the	290
236	sorting of the top $k-1$ elements might change in that cell (i.e., the first might become the second largest, and	291
237	vice versa), the coordinates that participate in the top $k-1$ support will be the same, while at the same time	292
238	the k -th largest will be $[\mathbf{v}(\varphi)]_i$.	293
239	Hence, for each coordinate $[\mathbf{v}(\varphi)]_i$, we need to visit the cells wherein it is the k -th largest. We do this by	294
240	examining <i>all</i> cells wherein $[\mathbf{v}(\varphi)]_i$ retains a fixed ranking. Visiting all these cells (T for each coordinate), is	295
241	possible by visiting all $n \cdot \binom{n-1}{d-1}$ intersection points of $\mathcal{B}_{i,j}(\varphi)$ curves as defined earlier. Since we know that each	296
242	cell is adjacent to at least 1 intersection point, then at each of these points we visit all adjacent cells. For each	297
243	cell that we visit, we compute the support of the largest k coordinates of a vector $\mathbf{v}(\varphi_0)$ with a φ_0 that lies in	298
244	that cell. We include this top k index set in \mathcal{S}_d and carry the same procedure for all cells. Since we visit all	299
245	coordinates and all their adjacent cells, this means that we visit all cells \mathcal{C}_j^i . This means that this procedure will	300
246	construct all possible supports in	301
247		302
248	$\mathcal{S}_d = \{\text{top}_k(c_1 \cdot \mathbf{v}_1 + \dots + c_d \cdot \mathbf{v}_d) : c_1, \dots, c_d \in \mathbb{R}\}$	303
249		304
250		305
251	1.3. Constructing the set \mathcal{S}_d	306
252	To visit all possible cells \mathcal{C}_j^i , we now have to check the intersection points, which are obtained by solving the	307
253	system of $d-1$ equations	308
254		309
255	$[\mathbf{v}(\varphi)]_{i_1} = [\mathbf{v}(\varphi)]_{i_2} = \dots = [\mathbf{v}(\varphi)]_{i_d}$	310
256	$\Leftrightarrow [\mathbf{v}(\varphi)]_{i_1} = [\mathbf{v}(\varphi)]_{i_2}, \dots, [\mathbf{v}(\varphi)]_{i_1} = [\mathbf{v}(\varphi)]_{i_d}.$	311
257		312
258	We can rewrite the above as	313
259	$\begin{bmatrix} \mathbf{e}_{i_1}^T - \mathbf{e}_{i_2}^T \\ \vdots \\ \mathbf{e}_{i_1}^T - \mathbf{e}_{i_d}^T \end{bmatrix} \mathbf{V}_d \mathbf{c} = \mathbf{0}_{(d-1) \times 1}$	314
260		315
261		316
262		317
263	where the solution is the nullspace of the matrix, which has dimension 1.	318
264	To explore all possible candidate vectors, we need to visit all cells. To do so, we compute all possible $\binom{n}{d}$ solution	319
265	intersection vectors $\mathbf{c}_{i_1, \dots, i_d}$. On each intersection vector we need to compute the locally optimal support set	320
266		321
267	$\text{top}_k(\mathbf{V}_d \mathbf{c}_{i_1, \dots, i_d}).$	322
268		323
269	Then observe that the coordinates i_1, \dots, i_d of $\mathbf{V}_d \mathbf{c}_{i_1, \dots, i_d}$ have the same value, since they all satisfy equation	324
270	(5). Let us assume that t of them appear in the set $\text{top}_k(\mathbf{V}_d \mathbf{c}_{i_1, \dots, i_d})$. Then, finding the top k supports of all	325
271	neighboring cell is equivalent to checking all different supports that can be generated by taking all $\binom{d}{t}$ possible t -	326
272	subsets of the i_1, \dots, i_d coordinates with respect to $\mathbf{V}_d \mathbf{c}_{i_1, \dots, i_d}$, while keeping the rest of the elements in $\mathbf{V}_d \mathbf{c}_{i_1, \dots, i_d}$	327
273	in their original ranking, as computed in $\text{top}_k(\mathbf{V}_d \mathbf{c}_{i_1, \dots, i_d})$. This, induces at most $O(\binom{d}{d/2})$ local sortings, i.e.,	328
274		329

330 top k supports. All these sortings will eventually be the elements of the \mathcal{S}_d set. The number of all candidate
 331 support sets will now be $O(\binom{d}{d/2} \binom{n-1}{d}) = O(n^d)$ and the total computation complexity is $O(n^{d+1})$, since for
 332 each point we compute the top- k support in linear time $O(n)$.
 333

334 For completeness the algorithm of the spannogram framework that generates \mathcal{S}_d is given below.
 335
 336

Algorithm 1 Spannogram Algorithm for \mathcal{S}_d .

```

338 1:  $\mathcal{S}_d = \emptyset$ 
339 2: for all  $(i_1, \dots, i_d) \in \{1, \dots, n\}^d$  and  $s \in \{-1, 1\}$  do
340 3:    $\mathbf{c} = s \cdot \text{nullspace} \left( \begin{bmatrix} [\mathbf{V}_d]_{i_1,:} & -[\mathbf{V}_d]_{i_2,:} \\ \vdots & \\ [\mathbf{V}_d]_{i_1,:} & -[\mathbf{V}_d]_{i_d,:} \end{bmatrix} \right)$ 
341 4:    $\mathbf{v} = \mathbf{V}_d^T \mathbf{c}$ 
342 5:    $\mathcal{S} = \text{top}_k(\mathbf{v})$ 
343 6:    $\mathcal{T} = \mathcal{S} - \{i_1, \dots, i_d\}$ 
344 7:   for all  $\binom{d}{d-|\mathcal{T}|}$  subsets  $\mathcal{J}$  of  $(i_1, \dots, i_d)$  do
345 8:      $\mathcal{S}_d = \mathcal{S}_d \cup (\mathcal{T} \cup \mathcal{J})$ 
346 9:   end for
347 10: end for
348 11: Output:  $\mathcal{S}_d$ .
```

1.4. Resolution of singularities

In our proofs, we assumed that the curves in $\mathbf{v}(\phi)$ are in general position. This is needed so that no more than $d-1$ curves intersect at a single point. This assumption is equivalent to requiring that every $d \times d$ submatrix of \mathbf{V}_d is full rank. This “general position” requirement can be handled by introducing infinitesimal perturbations in V_d . The details of the analysis of this method can be found in (Papailiopoulos et al., 2013).

2. Going from DkS to DBkS and back

In this subsection we show how a ρ -approximation algorithm for DBkS for arbitrary matrices, implies a 2ρ -approximation for DkS. Our proof goes through a randomized sampling argument.

At the end of this section we also show a deterministic scheme that converts a DBkS algorithm to an algorithm for DkS. This deterministic translation only guarantees that a ρ -approximation algorithm for DBkS for arbitrary matrices, can be converted to 4ρ -approximation for DkS. Our deterministic method uses much simpler vertex pruning techniques.

Algorithm 2 randombipartite(\mathcal{G})

```

372 1:  $\mathcal{L} = \emptyset, \mathcal{R} = \emptyset$ 
373 2: for Each vertex  $v$  in  $\mathcal{G}$  do
374 3:    $Z = \text{Bernoulli}(1/2)$ .
375 4:   if  $Z==1$  then
376 5:     Put  $v$  in  $\mathcal{L}$ 
377 6:   else
378 7:     Put  $v$  in  $\mathcal{R}$ 
379 8:   end if
380 9: end for
381 10:  $\mathcal{G}_B = \mathcal{G}$ 
382 11: delete all edges in  $\mathcal{G}_B(\mathcal{L})$  and  $\mathcal{G}_B(\mathcal{R})$ 
383 12: Output:  $\mathcal{G}_B$ 
```

440	2.1. Proof of Lemma 1: Randomized Reduction	495
441	Let us denote by $\mathcal{G}(\mathcal{S})$ the subgraph in \mathcal{G} induced by a vertex set \mathcal{S} . Let the adjacency matrix of the bipartite	496
442	graph created by <code>randombipartite</code> (\mathcal{G}) be \mathbf{G}_B	497
443		498
444	$\mathbf{A}_B = \begin{bmatrix} \mathbf{0}_{n_1 \times n_2} & \mathbf{B} \\ \mathbf{B}^T & \mathbf{0}_{n_2 \times n_1} \end{bmatrix},$	499
445		500
446		501
447	where $n_1 + n_2 = n$. In the following, we refer to \mathbf{B} as the bi-adjacency matrix of the bipartite graph \mathcal{G}_B . Moreover,	502
448	we denote as \mathcal{L} and \mathcal{R} the two disjoint vertex sets of a bipartite graph.	503
449		504
450	Before we proceed let us state a simple property on the quadratic form of bipartite graphs.	505
451	Proposition 1. Let $\mathbf{A}_B = \begin{bmatrix} \mathbf{0}_{n_1 \times n_2} & \mathbf{B} \\ \mathbf{B}^T & \mathbf{0}_{n_2 \times n_1} \end{bmatrix}$ be the adjacency matrix of a bipartite graph. Then, for any subset	506
452	of vertices \mathcal{S} , we have that $\mathcal{S} = \mathcal{S}_l \cup \mathcal{S}_r$, with $\mathcal{S}_l = \mathcal{S} \cap \mathcal{L}$ and $\mathcal{S}_r = \mathcal{S} \cap \mathcal{R}$. Moreover,	507
453		508
454	$\mathbf{1}_S^T \mathbf{A}_B \mathbf{1}_S = 2 \cdot \mathbf{1}_{\mathcal{S}_l}^T \mathbf{B} \mathbf{1}_{\mathcal{S}_r}$	509
455		510
456	<i>Proof.</i> It is easy to see that \mathcal{S}_l and \mathcal{S}_r are the vertex subsets of \mathcal{S} that correspond to either the left or right nodes	511
457	of the bipartite graph. Since the two sets are disjoint, we have	512
458		513
459	$\mathbf{1}_S = \mathbf{1}_{\mathcal{S}_l} + \mathbf{1}_{\mathcal{S}_r}.$	514
460		515
461	Then, the quadratic forms on \mathbf{A}_B can be equivalently rewritten as bilinear forms on \mathbf{B} :	516
462		517
463	$\mathbf{1}_S^T \mathbf{A}_B \mathbf{1}_S = (\mathbf{1}_{\mathcal{S}_l} + \mathbf{1}_{\mathcal{S}_r})^T \begin{bmatrix} \mathbf{0}_{n_1 \times n_2} & \mathbf{B} \\ \mathbf{B}^T & \mathbf{0}_{n_2 \times n_1} \end{bmatrix} (\mathbf{1}_{\mathcal{S}_l} + \mathbf{1}_{\mathcal{S}_r}) = \mathbf{1}_{\mathcal{S}_r}^T \mathbf{B}^T \mathbf{1}_{\mathcal{S}_l} + \mathbf{1}_{\mathcal{S}_l}^T \mathbf{B} \mathbf{1}_{\mathcal{S}_r} = 2 \cdot \mathbf{1}_{\mathcal{S}_l}^T \mathbf{B} \mathbf{1}_{\mathcal{S}_r}.$	518
464		519
465		520
466		521
467	Due to the above, we will consider the following DBkS problem	522
468		523
469	$\{\mathcal{X}_B, \mathcal{Y}_B\} = \arg \max_{ \mathcal{X} =k_1} \max_{ \mathcal{Y} =k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}},$	524
470		525
471		526
472	where $k_1 + k_2 = k$. Due to Proposition 1, $\mathcal{X}_B \cap \mathcal{Y}_B = \emptyset$, since the columns and rows of \mathbf{B} index two disjoint	527
473	vertex sets \mathcal{L} and \mathcal{R} , respectively. Then, our approximate solution with respect to the original DkS problem on	528
474	\mathbf{A} will then be	529
475	$\mathcal{S}_B = \mathcal{X}_B \cup \mathcal{Y}_B.$	530
476	Clearly $ \mathcal{S}_B = k$.	531
477		532
478	Proposition 2. The density of the above approximate solution is	533
479		534
480	$\text{den}(\mathcal{S}_B) = \frac{\mathbf{1}_{\mathcal{S}_B}^T \mathbf{A} \mathbf{1}_{\mathcal{S}_B}}{k} \geq \frac{\mathbf{1}_{\mathcal{S}_B}^T \mathbf{A}_B \mathbf{1}_{\mathcal{S}_B}}{k} = 2 \cdot \frac{\mathbf{1}_{\mathcal{X}_B}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}_B}}{k}.$	535
481		536
482		537
483	<i>Proof.</i> The result follows immediately by the nonnegativity of the entries in \mathbf{A} , and the fact that \mathbf{A}_B contains a	538
484	subset of the entries of \mathbf{A} . The last equality follows from Proposition 1. \square	539
485		540
486	We will now show that $\text{den}(\mathcal{S}_B)$ is at least $\text{opt}/(2 + \delta)$ with high probability, if we solve DBkS on $\log n$ graphs	541
487	independently created using <code>randombipartite</code> (\mathcal{G}), and by keeping the best solution among all $\log n$ extracted	542
488	sets \mathcal{S}_B .	543
489	Proposition 3. Let us partition the vertices of \mathcal{G} into two sets \mathcal{L}, \mathcal{R} according to <code>randombipartite</code> (\mathcal{G}): We	544
490	flip a fair coin for each vertex of \mathcal{G} and put it with probability $1/2$ in either of the two sets. Then, we create a	545
491	bipartite graph \mathcal{G}_B that has as left and right vertex sets the sets \mathcal{L} and \mathcal{R} , respectively. The edges that we maintain	546
492	from the original graph, are only those that cross from \mathcal{L} to \mathcal{R} , while we delete all edges in the subgraphs induced	547
493	by both sets \mathcal{L} and \mathcal{R} . Then, there exists in a k -subgraph in \mathcal{G}_B , that contains $0.5 \cdot k \cdot \text{opt}$ edges, in expectation.	548
494		549

550 *Proof.* First observe that we can represent the edges of \mathcal{G}_B as random variables $X_{i,j}$. If (i,j) is not an edge in
 551 \mathcal{G} , then $X_{i,j}$ will be 0 with probability 1. If however (i,j) is an edge in \mathcal{G} , then $X_{i,j}$ is 1, i.e., appears in \mathcal{G}_B ,
 552 with the same probability that one of its vertices lands in \mathcal{L} , while the second is in \mathcal{R} . It is easy to find that this
 553 probability is $\Pr\{X_{i,j} = 1\} = 1/2$. Hence,
 554

$$555 \quad X_{i,j} = \begin{cases} 0, & \text{if } (i,j) \text{ not an edge in } \mathcal{G}, \\ 556 \quad Z, & \text{if } (i,j) \text{ is an edge in } \mathcal{G}, \end{cases} \quad 609$$

558 where Z is a Bernoulli($1/2$) random variable.
 559

560 Now let \mathcal{S}_* denote the vertex set of the densest k -subgraph on the original graph \mathcal{G} , that has density $\text{den}(\mathcal{S}_*) = \text{opt}$.
 561 Observe that for that subgraph we have
 562

$$562 \quad \mathbf{1}_{\mathcal{S}_*}^T \mathbf{A} \mathbf{1}_{\mathcal{S}_*} = \sum_{i,j \in \mathcal{S}_*} A_{i,j} = k \cdot \text{opt}. \quad 617$$

565 Let \mathbf{A}_B denote the adjacency matrix of the bipartite graph \mathcal{G}_B . Then, we have that the expected quadratic form
 566 on the new adjacency $\mathbf{1}_{\mathcal{S}_*}^T \mathbf{A}_B \mathbf{1}_{\mathcal{S}_*}$ is:
 567

$$568 \quad E \left\{ \mathbf{1}_{\mathcal{S}_*}^T \mathbf{A}_B \mathbf{1}_{\mathcal{S}_*} \right\} = E \left\{ \sum_{i,j \in \mathcal{S}_*} X_{i,j} \right\} = E \left\{ \sum_{i,j \in \mathcal{S}_* \text{ and } (i,j) \in \mathcal{G}} X_{i,j} \right\} = E \left\{ \sum_{i,j \in \mathcal{S}_* \text{ and } (i,j) \in \mathcal{G}} Z \right\} \quad 623$$

$$569 \quad = \frac{1}{2} \cdot \sum_{i,j \in \mathcal{S}_*} A_{i,j} = 0.5 \cdot k \cdot \text{opt}. \quad 624$$

574 \square
 575

577 We will now show that if we run `randombipartite`(\mathcal{G}) a total number of $3 \log n \cdot \log \log n$ times, then with high
 578 probability, at least one \mathcal{G}_B will contain a k -subgraph with density at least $0.5 \cdot \text{opt}$. This will imply that the
 579 densest k subgraph of \mathcal{G}_B will have density *at least* $0.5 \cdot \text{opt}$.
 580

581 **Proposition 4.** *Let us run `randombipartite`(\mathcal{G}) $m = 3 \log \cdot \log \log n$ times, and let us obtain each time a graph
 582 \mathcal{G}_B^i . For each bipartite graph let \mathbf{A}_B^i denote its adjacency and \mathbf{B}_i its bi-adjacency, let*
 583

$$584 \quad \{\mathcal{X}_i, \mathcal{Y}_i\} = \arg \max_{k_1+k_2=k} \max_{|\mathcal{X}|=k_1} \max_{|\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B}_i \mathbf{1}_{\mathcal{Y}}, \quad 638$$

586 and let $\mathcal{S}_i = \mathcal{X}_i \cup \mathcal{Y}_i$ denote the optimal k -subgraph of \mathcal{G}_B^i . Then, with probability at least $1 - \frac{1}{\log n}$, we have
 587

$$588 \quad \max_{i=1,\dots,m} \text{den}(\mathcal{S}_i) = \max_{i=1,\dots,m} \mathbf{1}_{\mathcal{S}_i}^T \mathbf{A} \mathbf{1}_{\mathcal{S}_i} \geq \frac{\text{opt}}{2 + 1/\log n}. \quad 643$$

591 *Proof.* In this proof we will use the reverse Markov Inequality which states that for any random variable X , such
 592 that $X \leq m$, then, for any $a \leq E\{X\}$, we have
 593

$$594 \quad \Pr\{X \leq a\} \leq \frac{m - E\{X\}}{m - a}. \quad 649$$

597 Let \mathcal{S}_* denote the densest k -subgraph for \mathcal{G} . Here, our random variable will be the the quadratic form
 598

$$599 \quad X = \mathbf{1}_{\mathcal{S}_*}^T \mathbf{A}_B^i \mathbf{1}_{\mathcal{S}_*}. \quad 654$$

601 Due to Proposition 3, we have that $E\{X\} = 0.5 \cdot k \cdot \text{opt}$. Moreover, we know that
 602

$$603 \quad \mathbf{1}_{\mathcal{S}_*}^T \mathbf{A}_B^i \mathbf{1}_{\mathcal{S}_*} \leq \mathbf{1}_{\mathcal{S}_*}^T \mathbf{A} \mathbf{1}_{\mathcal{S}_*} = k \cdot \text{opt}. \quad 658$$

604

660 Hence, we will set $\delta = 1/\log n$ and use the reverse Markov inequality:

$$\begin{aligned}
 \Pr \left\{ \mathbf{1}_{S_*}^T \mathbf{A}_B^i \mathbf{1}_{S_*} \leq k \cdot \text{opt}/(2 + 1/\log n) \right\} &\leq \frac{k \cdot \text{opt} - k \cdot \text{opt}/2}{k \cdot \text{opt} - k \cdot \text{opt}/(2 + 1/\log n)} \\
 &= \frac{1 - 1/2}{1 - 1/(2 + 1/\log n)} = \frac{1/2}{\frac{1+1/\log n}{2+1/\log n}} \\
 &= \frac{2 + 1/\log n}{2 + 2/\log n} = \frac{2 \log n + 1}{2 \log n + 2} = 1 - \frac{1}{2(\log n + 1)} \leq 1 - \frac{1}{3 \log n}.
 \end{aligned}$$

670 Now, let B_i denote the event that $\{\mathbf{1}_{S_*}^T \mathbf{A}_B^i \mathbf{1}_{S_*} \leq k \cdot \text{opt}/(2 + 1/\log n)\}$. We would like to find a number of draws
 671 m such that, with high probability $1 - \frac{1}{\log n}$, we will have one graph \mathcal{G}_B^i satisfying
 672

$$\mathbf{1}_{S_*}^T \mathbf{A}_B^i \mathbf{1}_{S_*} > k \cdot \text{opt}/(2 + 1/\log n).$$

673 Since these m events are independent for different graphs, it is easy to see find the number of draws m by solving
 674 the equation:
 675

$$\left(1 - \frac{1}{3 \log n}\right)^m \leq \frac{1}{\log n} \Rightarrow m \cdot \log \left(1 - \frac{1}{3 \log n}\right) \leq -\log(\log n) \Rightarrow m \geq \frac{-\log(\log n)}{\log \left(1 - \frac{1}{3 \log n}\right)} \Rightarrow m \geq 3 \log n \cdot \log \log n.$$

681 where in the last step we used the fact that $-\log(1 - \epsilon) > \epsilon$, for $\epsilon < 1$.

682 The above derivations tell us that if we draw $m = 3 \log n \cdot \log \log n$ graphs \mathcal{G}_B^i using `randombipartite(G)`, then
 683 at least one of them contains a subgraph of density
 684

$$\frac{\text{opt}}{2 + 1/\log n}.$$

685 Hence, the subgraph $S_i = \mathcal{X}_i \cup \mathcal{Y}_i$ with the highest density among the m densest k -subgraph of the graphs \mathcal{G}_B^i
 686 has density at least $\frac{\text{opt}}{2+1/\log n}$. \square

687 Using Proposition 4, we can now establish Lemma 1 in the following way. Suppose that we have an approximation
 688 algorithm such that for any k_1, k_2 , we can obtain a subgraph with density \mathcal{A}

$$\frac{1}{\rho} \cdot \arg \max_{|\mathcal{X}|=k_1} \max_{|\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}} \leq \mathcal{A} \leq \arg \max_{|\mathcal{X}|=k_1} \max_{|\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}}$$

689 Then, we could use it with Proposition 4, to obtain a $(2+\delta)\rho$ approximation algorithm for DkS, for any $\delta > 1/n^\gamma$,
 690 and $\gamma = \Theta(1)$, in polynomial time.
 691

700 What we show in our approximation results is that we can compute a solution with density at least
 701

$$\arg \max_{|\mathcal{X}|=k_1} \max_{|\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}} - 2|\lambda_{d+1}|,$$

704 where λ_{d+1} is the $d+1$ absolutely largest eigenvalue of the adjacency matrix \mathbf{A} .
 705

2.2. Deterministic Reduction

708 Here, we show a simper translation of a DBkS algorithm to a DkS algorithm.

709 **Proposition 5.** Let \mathcal{X}, \mathcal{Y} be two vertex sets of size k , and let $\mathcal{Z} = \mathcal{X} \cup \mathcal{Y}$. Then, we can use `prune(G, X*, Y*, k)` to construct in time $O(k^2)$ a vertex set \mathcal{Z}' of size k , that has density at least
 710

$$\left(\frac{1}{4} - O\left(\frac{1}{k}\right)\right) \text{den}(\mathcal{Z}).$$

770	Algorithm 3 <code>prune</code> ($\mathcal{G}, \mathcal{X}, \mathcal{Y}, k$)	825
771	1: Input: $\mathcal{G}, \mathcal{X}, \mathcal{Y}$	826
772	2: $\mathcal{Z}_1 = \mathcal{X} \cup \mathcal{Y}, k' = \mathcal{Z}_1 $	827
773	3: for $i = 1 : k' - k$ do	828
774	4: find v smallest degree vertex in $\mathcal{G}(\mathcal{Z}_1)$	829
775	5: $\mathcal{Z}_{i+1} = \mathcal{Z}_i \setminus \{v\}$.	830
776	6: end for	831
777	7: Output: $\mathcal{Z} = \mathcal{Z}_{k'-k}$	832
778		833
779		834

Proof. Let k' denote the size of the set $\mathcal{Z}_1 = \mathcal{X} \cup \mathcal{Y}$. Clearly $k \leq k' \leq 2k$. Let us denote as $d_{\min}(\mathcal{S})$, the smallest degree in the subgraph induced by the vertex set \mathcal{S} . Then, observe that its minimum degree can be bounded as

$$d_{\min}(\mathcal{S}) \leq \frac{\text{den}(\mathcal{S})}{|\mathcal{S}|}.$$

During the first iteration of `prune`, we delete the smallest degree vertex from \mathcal{Z}_1 and obtain \mathcal{Z}_2 . This leads to a loss in density that is equal to $2 \cdot d_{\min}(\mathcal{Z}_1)$. Hence, the density of \mathcal{Z}_2 is lower bounded as

$$\text{den}(\mathcal{Z}_2) = \text{den}(\mathcal{Z}_1) - 2 \cdot d_{\min}(\mathcal{Z}_1) \geq \text{den}(\mathcal{Z}_1) - 2 \cdot \frac{\text{den}(\mathcal{Z}_1)}{k} = \frac{k-2}{k} \cdot \text{den}(\mathcal{Z}_1).$$

This means that for any iteration i of the algorithm the density of the current subgraph induced by \mathcal{Z}_i is at least

$$\text{den}(\mathcal{Z}_{i+1}) \geq \frac{(k' - i + 1) - 2}{k' - i + 1} \cdot \text{den}(\mathcal{Z}_i) = \frac{k - i - 1}{k' - i + 1} \cdot \text{den}(\mathcal{Z}_i).$$

Hence, the output vertex set \mathcal{Z} satisfies

$$\text{den}(\mathcal{Z}) \geq \left(\prod_{i=1}^{k-k} \frac{k' - i - 1}{k' - i + 1} \right) \text{den}(\mathcal{Z}_1).$$

We can now compute the product using factorials. For the numerator we have

$$(k' - 2) \cdot (k' - 3) \cdot \dots \cdot (k' - l - 1) = \frac{k'!}{k' \cdot (k' - 1) \cdot [(k' - l - 2)!]}.$$

And for the denominator

$$k' \cdot (k' - 1) \cdot \dots \cdot (k' - l + 1) = \frac{k'!}{(k' - l)!}.$$

Hence,

$$\prod_{i=1}^l \frac{k' - i - 1}{k' - i + 1} = \frac{(k' - l)!}{k' \cdot (k' - 1) \cdot [(k' - l - 2)!]} = \frac{(k' - l - 1)(k' - l)}{k'(k' - 1)}.$$

Hence, for $l = k' - k$ we get

$$\text{den}(\mathcal{Z}) = \frac{k \cdot (k - 1)}{k' \cdot (k' - 1)} \geq \frac{k \cdot (k - 1)}{2k \cdot (2k - 1)} = \frac{1}{4} - \frac{1}{8k - 2}$$

which establishes the result. \square

Hence, if we had an algorithm that obtains an approximate solution $\{\hat{\mathcal{X}}, \hat{\mathcal{Y}}\}$ such that

$$\mathbf{1}_{\hat{\mathcal{X}}}^T \mathbf{A} \mathbf{1}_{\hat{\mathcal{Y}}} \geq \frac{\mathbf{1}_{\mathcal{X}^*}^T \mathbf{A} \mathbf{1}_{\hat{\mathcal{Y}}_*}}{\rho}.$$

Then, we could use `prune` to obtain a k -subgraph $\hat{\mathcal{Z}}$, such that

$$\text{den}(\hat{\mathcal{Z}}) \geq \frac{1}{4} \frac{\mathbf{1}_{\hat{\mathcal{X}}}^T \mathbf{A} \mathbf{1}_{\hat{\mathcal{Y}}}}{k} \geq \frac{1}{4 \cdot \rho} \frac{\mathbf{1}_{\mathcal{X}^*}^T \mathbf{A} \mathbf{1}_{\hat{\mathcal{Y}}_*}}{k} \geq \frac{1}{4 \cdot \rho} \frac{\mathbf{1}_{\mathcal{S}^*}^T \mathbf{A} \mathbf{1}_{\hat{\mathcal{S}}_*}}{k} \geq \frac{1}{4 \cdot \rho} \text{opt}.$$

3. Approximation Guarantees
3.1. Proof of Theorem 1

The first important technical proposition that we show, is that we can solve DBkS for any constant rank *rectangular* matrix \mathbf{B} of dimensions $n_1 \times n_2$.

Proposition 6. *Let \mathbf{B} be any matrix of size $n_1 \times n_2$ and let*

$$\mathbf{B}_d = \sum_{i=1}^d \sigma_i \mathbf{v}_i \mathbf{u}_i^T$$

be its singular value decomposition, where \mathbf{v}_i and \mathbf{u}_i is the left and right singular vectors corresponding to the i th largest singular value $\sigma_i(\mathbf{B})$. Then, we can solve the following problem

$$\{\mathcal{X}_d, \mathcal{Y}_d\} = \arg \max_{|\mathcal{X}|=k_1, |\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B}_d \mathbf{1}_{\mathcal{Y}}.$$

in time $O(\min\{n_1, n_2\}^{d+1})$.

Proof. For simplicity we assume that the left singular vectors are scaled by their singular values, hence

$$\mathbf{B}_d = \mathbf{v}_1 \mathbf{u}_1^T + \dots + \mathbf{v}_d \mathbf{u}_d^T.$$

Let us without loss of generality assume that $n_1 \leq n_2$.

We wish to solve:

$$\max_{|\mathcal{X}|=k_1, |\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T (\mathbf{v}_1 \mathbf{u}_1^T + \dots + \mathbf{v}_d \mathbf{u}_d^T) \mathbf{1}_{\mathcal{Y}}. \quad (8)$$

Observe that we can rewrite (8) in the following way

$$\max_{|\mathcal{X}|=k_1, |\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T \left[\mathbf{v}_1 \cdot \underbrace{(\mathbf{u}_1^T \mathbf{1}_{\mathcal{Y}})}_{c_1} + \dots + \mathbf{v}_d \cdot \underbrace{(\mathbf{u}_d^T \mathbf{1}_{\mathcal{Y}})}_{c_d} \right] = \max_{|\mathcal{Y}|=k_2} \left(\max_{|\mathcal{X}|=k_1} \mathbf{1}_{\mathcal{X}}^T \mathbf{v}_{\mathcal{Y}} \right),$$

where $\mathbf{v}_{\mathcal{Y}} = \mathbf{v}_1 \cdot c_1 + \dots + \mathbf{v}_d \cdot c_d$ is an n_1 -dimensional vector generated by the d -dimensional subspace spanned by $\mathbf{v}_1, \dots, \mathbf{v}_d$.

We will now make a key observation: for every fixed vector $\mathbf{v}_{\mathcal{Y}}$, the index set \mathcal{X} that maximizes $\mathbf{1}_{\mathcal{X}}^T \mathbf{v}_{\mathcal{Y}}$ can be easily computed. It is not hard to see that for any fixed vector $\mathbf{v}_{\mathcal{Y}}$, the k_1 -subset \mathcal{X} that maximizes

$$\mathbf{1}_{\mathcal{X}}^T \mathbf{v}_{\mathcal{Y}} = \sum_{i \in \mathcal{X}} [\mathbf{v}_{\mathcal{Y}}]_i$$

corresponds to either the set of k_1 largest or k_1 smallest signed coordinates of $\mathbf{v}_{\mathcal{Y}}$. That is, the locally optimal sets are either $\text{top}_{k_1}(\mathbf{v}_{\mathcal{Y}})$ or $\text{top}_{k_1}(-\mathbf{v}_{\mathcal{Y}})$.

We now wish to find all possible locally optimal sets \mathcal{X} . If we could possibly check all vectors $\mathbf{v}_{\mathcal{Y}}$, then we could find all locally optimal index sets $\text{top}_{k_1}(\pm \mathbf{v}_{\mathcal{Y}})$.

Let us denote as \mathcal{S}_d the set of all k_1 -sized sets \mathcal{X} that are the optimal solutions of the inner maximization of in the above, for *any* vector \mathbf{v} in the span of $\mathbf{v}_1, \dots, \mathbf{v}_d$

$$\mathcal{S}_d = \{\text{top}_{k_1}(\pm [\mathbf{v}_1 \cdot c_1 + \dots + \mathbf{v}_d \cdot c_d]) : c_1, \dots, c_d \in \mathbb{R}\}.$$

Clearly, this set contains all possible locally optimal \mathcal{X} sets of the form $\text{top}_{k_1}(\mathbf{v}_{\mathcal{Y}})$. Therefore, we can rewrite DBkS on \mathbf{B}_d as

$$\max_{|\mathcal{Y}|=k_2} \max_{\mathcal{X} \in \mathcal{S}_d} \mathbf{1}_{\mathcal{X}}^T \mathbf{B}_d \mathbf{1}_{\mathcal{Y}}. \quad (9)$$

- 990 The above problem can now be solved in the following way: for every set $\mathcal{X} \in \mathcal{S}_d$ find the locally optimal set \mathcal{Y}
 991 that maximizes $\mathbf{1}_{\mathcal{X}}^T \mathbf{B}_d \mathbf{1}_{\mathcal{Y}}$. Again, this will either be $\text{top}_{k_2}(-\mathbf{B}_d \mathbf{1}_{\mathcal{X}})$ or $\text{top}_{k_2}(\mathbf{B}_d \mathbf{1}_{\mathcal{X}})$. Then, we simply need to
 992 test all such \mathcal{X}, \mathcal{Y} pairs on \mathbf{B}_d and keep the optimizer.
 993 Due to the above, the problem of solving DBkS on the rectangular matrix \mathbf{B}_d is equivalent to constructing the set
 994 of k_1 -supports \mathcal{S}_d , and then finding the optimal solution in that set. How large can \mathcal{S}_d be and can we construct
 995 it in polynomial time? As we showed in the first section of the supplemental material this set has size $O(\binom{n_1}{d})$
 996 and can be constructed in time $O(n_1^{d+1})$.
 997
 998 Observe that in the above we could have equivalently solved the problem by finding all the top k_2 sets in the
 999 span of $\mathbf{u}_1, \dots, \mathbf{u}_d$, say that they belong in set \mathcal{S}'_d . Then, we could solve the problem by finding for each k_2 sized
 1000 set $\mathcal{Y} \in \mathcal{S}'_d$ the optimal k_1 sized set \mathcal{X} . Both approaches are the same, and the one with the smallest dimension
 1001 is selected to reduce the computational complexity.
 1002
 1003 The algorithm that solves the problem for rectangular matrices is given below.
 1004
 1005 **Algorithm 4** low-rank approximations for DBkS
 1006 1: **lowrankDBkS**(k_1, k_2, d, \mathbf{B})
 1007 2: $[\mathbf{V}_d, \Sigma_d, \mathbf{U}_d] = \text{SVD}(\mathbf{B}, d)$
 1008 3: $\mathcal{S}_d = \text{Spannogram}(k_1, \mathbf{V}_d)$
 1009 4: $\{\mathcal{X}_d, \mathcal{Y}_d\} = \arg \max_{|\mathcal{Y}|=k_2} \max_{\mathcal{X} \in \mathcal{S}_d} \mathbf{1}_{\mathcal{X}}^T \mathbf{V}_d \Sigma_d \mathbf{U}_d^T \mathbf{1}_{\mathcal{Y}}$
 1010 5: **Output:** $\{\mathcal{X}_d, \mathcal{Y}_d\}$
 1011
 1012 1: **Spannogram**(k_1, \mathbf{V}_d)
 1013 2: $\mathcal{S}_d = \{\text{top}_k(\mathbf{v}) : \mathbf{v} \in \text{span}(\mathbf{v}_1, \dots, \mathbf{v}_d)\}$
 1014 3: **Output:** \mathcal{S}_d .
 1015
 1016
 1017
 1018
 1019 In our following derivations, for both cases of a rectangular and square symmetric matrices, we consider the same
 1020 notation of the output solution and output density for simplicity:
 1021
 1022
 1023 $\{\mathcal{X}_d, \mathcal{Y}_d\} = \arg \max_{|\mathcal{X}|=k, |\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}}$ and $\text{opt}_d^{\mathbf{A}} = \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}_d}}{k}$,
 1024
 1025 $\{\mathcal{X}_d, \mathcal{Y}_d\} = \arg \max_{k_1, k_2: k_1+k_2=k} \max_{|\mathcal{X}|=k_1, |\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B}_d \mathbf{1}_{\mathcal{Y}}$ and $\text{opt}_d^{\mathbf{B}} = 2 \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}_d}}{k}$.
 1026
 1027 Moreover, as a reminder the optimal solutions and densities for the problems of interest (DkS, DBkS on \mathbf{A} , and
 1028 DBkS on \mathbf{B}) are
 1029
 1030
 1031 $\mathcal{S}^* = \arg \max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A} \mathbf{1}_{\mathcal{S}}$ and $\text{opt} = \frac{\mathbf{1}_{\mathcal{S}^*}^T \mathbf{A} \mathbf{1}_{\mathcal{S}^*}}{k}$,
 1032
 1033 $\{\mathcal{X}^*, \mathcal{Y}^*\} = \arg \max_{|\mathcal{X}|=k, |\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}}$ and $\text{opt}^{\mathbf{A}} = \frac{\mathbf{1}_{\mathcal{X}^*}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}^*}}{k}$,
 1034
 1035 $\{\mathcal{X}^*, \mathcal{Y}^*\} = \arg \max_{k_1, k_2: k_1+k_2=k} \max_{|\mathcal{X}|=k_1, |\mathcal{Y}|=k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}}$ and $\text{opt}^{\mathbf{B}} = 2 \frac{\mathbf{1}_{\mathcal{X}^*}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}^*}}{k}$.
 1036
 1037
 1038 We continue with bounding the distance between the optimal solution for DBkS and rank- d optimal solution pair
 1039 $\{\mathcal{X}_d, \mathcal{Y}_d\}$. We have the following result.
 1040
 1041 **Proposition 7.** For any matrix \mathbf{A} , we have
 1042
 1043
$$\text{opt}_d^{\mathbf{B}} \geq \text{opt}^{\mathbf{B}} - 2 \cdot |\lambda_{d+1}|. \quad (10)$$

 1044

1100 Moreover, for any rectangular matrix \mathbf{B} , we have

$$1101 \quad \text{opt}_d^{\mathbf{B}} \geq \text{opt}^{\mathbf{B}} - 2 \cdot \sigma_{d+1}. \quad (11)$$

1103 Proof. Let $\mathcal{X}_*, \mathcal{Y}_*$ be the optimal solution of DBkS on \mathbf{A} and let $\mathcal{X}_d, \mathcal{Y}_d$ be the optimal solution of DBkS on the
1104 rank- d matrix \mathbf{A}_d . Then, we have

$$\begin{aligned} 1106 \quad \text{opt}_d^{\mathbf{B}} &= \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}_d}}{k} = \frac{\mathbf{1}_{\mathcal{X}_d}^T (\mathbf{A}_d + \mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_d}}{k} = \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} + \frac{\mathbf{1}_{\mathcal{X}_d}^T (\mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_d}}{k} \\ 1107 &\geq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} - \frac{\|\mathbf{1}_{\mathcal{X}_d}\|_2 \cdot \|(\mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_d}\|_2}{k} \geq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} - |\lambda_{d+1}|, \end{aligned} \quad (12)$$

1111 where the first inequality comes due to Cauchy-Schwarz and the second due to the fact that the norm of the
1112 indicator vector is k and the operator norm of $\mathbf{A} - \mathbf{A}_d$ is equal to the $d+1$ largest eigenvalue of \mathbf{A} .

1113 Moreover, we have that

$$\begin{aligned} 1115 \quad \text{opt}^{\mathbf{B}} &= \frac{\mathbf{1}_{\mathcal{X}_*}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}_*}}{k} = \frac{\mathbf{1}_{\mathcal{X}_*}^T (\mathbf{A}_d + \mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_*}}{k} = \frac{\mathbf{1}_{\mathcal{X}_*}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_*}}{k} + \frac{\mathbf{1}_{\mathcal{X}_*}^T (\mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_*}}{k} \\ 1116 &\leq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} + \frac{\mathbf{1}_{\mathcal{X}_*}^T (\mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_*}}{k} \leq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} + \frac{\|\mathbf{1}_{\mathcal{X}_*}\|_2 \cdot \|(\mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_*}\|_2}{k} \leq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} + |\lambda_{d+1}| \end{aligned} \quad (13)$$

1120 where the first inequality comes due to the fact that $\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d} \geq \mathbf{1}_{\mathcal{X}_*}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_*}$ and the second and third are similar
1121 to the previous bound. We can now combine the above two bound to obtain:

$$1122 \quad \text{opt}_d^{\mathbf{B}} \geq \text{opt}^{\mathbf{B}} - 2 \cdot |\lambda_{d+1}|. \quad (14)$$

1124 In the exact same way, we can obtain the result for rectangular matrices. \square

1125 The above proposition, combined with Proposition 1 give us the bipartite part of Theorem 1, where $\text{opt}^{\mathbf{B}} = \text{opt}$,
1126 that is

$$1128 \quad \text{opt}_d^{\mathbf{B}} \geq \text{opt}^{\mathbf{B}} - 2 \cdot |\lambda_{d+1}| = \text{opt} - 2 \cdot |\lambda_{d+1}|.$$

1130 To establish the part about graphs with the d largest eigenvalues being positive, we use the following result.

1131 **Proposition 8.** If \mathbf{A}_d is positive semidefinite, then

$$1132 \quad \max_{|\mathcal{X}|=k} \frac{\mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{X}}}{k} = \max_{|\mathcal{X}|=k} \max_{|\mathcal{Y}|=k} \frac{\mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}}}{k}$$

1135 Proof. This is easy to see by the fact that for any two sets \mathcal{X}, \mathcal{Y} we have

$$\begin{aligned} 1137 \quad \max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}} &\leq \max_{|\mathcal{X}|=k, |\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}} = \max_{|\mathcal{X}|=k, |\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{V}_d \mathbf{\Lambda}_d^{1/2} \mathbf{\Lambda}_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{Y}} \\ 1138 &\leq \max_{|\mathcal{X}|=k, |\mathcal{Y}|=k} \max \left\{ \|\mathbf{\Lambda}_d^{1/2} \mathbf{V}_d \mathbf{1}_{\mathcal{X}}\|^2, \|\mathbf{\Lambda}_d^{1/2} \mathbf{V}_d \mathbf{1}_{\mathcal{Y}}\|^2 \right\} \leq \max_{|\mathcal{X}|=k, |\mathcal{Y}|=k} \max \{ \mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{X}}, \mathbf{1}_{\mathcal{Y}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}} \} \\ 1139 &\leq \max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}} \end{aligned}$$

1144 where the second inequality comes due to the Cauchy-Schwarz inequality. \square

1146 We can combine the above proposition with the first part of Proposition 7 to obtain that

$$1148 \quad \text{opt}_d^{\mathbf{B}} \geq \text{opt} - 2|\lambda_{d+1}(\mathbf{A})|$$

1150 when \mathbf{A}_d is positive semidefinite.

1151 In the next proposition, we show how to translate a low-rank approximation of \mathbf{A} after we used the random
1152 sampling of `randombipartite(G)`. We need this result to establish the general result of Theorem 1, by
1153 connecting the previous spectral bound, with the 2 loss in approximation between DBkS and DkS.
1154

1210 **Proposition 9.** Let \mathbf{A} be the adjacency matrix of a graph. Moreover, let the matrices \mathbf{P}_1 and \mathbf{P}_2 be such that
 1211 $\mathbf{B} = \mathbf{P}_1 \mathbf{A} \mathbf{P}_2$ is the bi-adjacency created by each loop of $\text{randombipartite}(\mathcal{G})$, where \mathbf{P}_1 is an $n_1 \times n$ matrix
 1212 indexing the left vertices of the graph, and \mathbf{P}_2 is an $n \times n_2$ sampling matrix that indexes the right vertices of the
 1213 sub-sampled graph. Then,

$$1215 \quad \text{opt}_d^{\mathbf{B}} \geq \text{opt}^{\mathbf{B}} - 2|\lambda_{d+1}(\mathbf{A})|,$$

1218 where $\text{opt}^{\mathbf{B}}$ is the maximum density on $\mathbf{B} = \mathbf{P}_1 \mathbf{A} \mathbf{P}_2$.

1223 *Proof.* Let without loss of generality assume that \mathbf{B} will be the bipartite subgraph between the first n_1 and the
 1224 remaining $n_2 = n - n_1$ vertices, such that

$$1226 \quad \mathbf{A} = \begin{bmatrix} \mathbf{C} & \mathbf{B} \\ \mathbf{B}^T & \mathbf{D} \end{bmatrix} \quad (15)$$

1229 Then there are two sampling matrices that pick the corresponding columns and rows

$$1232 \quad \mathbf{P}_1 = [\mathbf{I}_{n_1 \times n_1} \quad \mathbf{0}_{n_1 \times n_2}] \text{ and } \mathbf{P}_2 = \begin{bmatrix} \mathbf{0}_{n_1 \times n_2} \\ \mathbf{I}_{n_2 \times n_2} \end{bmatrix}$$

1235 Then, instead of working on the matrix that is the rank- d best fit for \mathbf{B} , we work on

$$1238 \quad \mathbf{B}_d = \mathbf{P}_1 \mathbf{A}_d \mathbf{P}_2.$$

1241 Now, we use the bounding techniques of our previous derivations:

$$1243 \quad \text{opt}_d = \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}_d}}{k} = \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{P}_1 (\mathbf{A}_d + \mathbf{A} - \mathbf{A}_d) \mathbf{P}_2 \mathbf{1}_{\mathcal{Y}_d}}{k} = \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{P}_1 \mathbf{A}_d \mathbf{P}_2 \mathbf{1}_{\mathcal{Y}_d}}{k} + \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{P}_1 (\mathbf{A} - \mathbf{A}_d) \mathbf{P}_2 \mathbf{1}_{\mathcal{Y}_d}}{k} \\ 1244 \geq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{P}_1 \mathbf{A}_d \mathbf{P}_2 \mathbf{1}_{\mathcal{Y}_d}}{k} - \frac{\|\mathbf{1}_{\mathcal{X}_d}\|_2 \cdot \|\mathbf{P}_1 (\mathbf{A} - \mathbf{A}_d) \mathbf{P}_2 \mathbf{1}_{\mathcal{Y}_d}\|_2}{k} \geq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{P}_1 \mathbf{A}_d \mathbf{P}_2 \mathbf{1}_{\mathcal{Y}_d}}{k} - |\lambda_{d+1}|, \quad (16)$$

1248 where the last step comes due to the fact that $\mathbf{P}_1, \mathbf{P}_2$ their singular values are 1. We can use a similar bound to
 1249 obtain

$$1252 \quad \text{opt}^{\mathbf{B}} \leq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{P}_1 \mathbf{A}_d \mathbf{P}_2 \mathbf{1}_{\mathcal{Y}_d}}{k} + |\lambda_{d+1}|,$$

1255 where $\text{opt}^{\mathbf{B}}$ is the density of the densest k -subgraph on the graph with bi-adjacency matrix $\mathbf{P}_1 \mathbf{A} \mathbf{P}_2$. and combine
 1256 the above to establish the result. \square

1261 We can now use our random sampling Proposition 4 and combine that with Propositions 9, and 7, to establish
 1262 Theorem 1 for arbitrary graphs.

1263 The general algorithm is given below

1320 **Algorithm 5** low-rank approximations for DkS 1375
 1321 1: $[\mathbf{V}_d, \mathbf{\Lambda}_d] = \text{EVD}(\mathbf{A}, d)$ 1376
 1322 2: **if** \mathcal{G} is bipartite **then** 1377
 1323 3: \mathbf{B} =bi-adjacency of \mathcal{G} 1378
 1324 4: $[\mathbf{V}_d, \mathbf{\Sigma}_d, \mathbf{U}_d] = \text{SVD}(\mathbf{B}, d)$ 1379
 1325 5: $\mathbf{B}_d = \mathbf{V}_d \mathbf{\Sigma}_d \mathbf{U}_d^T$ 1380
 1326 6: $\{\mathcal{X}_d, \mathcal{Y}_d\} = \arg \max_{|\mathcal{X}|+|\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{B}_d \mathbf{1}_{\mathcal{Y}}$. 1381
 1327 7: $\mathcal{Z}_d = \mathcal{X}_d \cup \mathcal{Y}_d$ 1382
 1328 8: **else if** \mathcal{G} is not bipartite and first d eigenvalues are positive **then** 1383
 1329 9: $\{\mathcal{X}_d, \mathcal{X}_d\} = \arg \max_{|\mathcal{X}|=k} \max_{|\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}}$. 1384
 1330 10: $\mathcal{Z}_d = \mathcal{X}_d$ 1385
 1331 11: **else if** \mathcal{G} is not bipartite and first d eigenvalues not all positive **then** 1386
 1332 12: **for** $i = 1 : O(\text{poly}(\log n))$ **do** 1387
 1333 draw n fair coins and assign them to vertices 1388
 1334 \mathcal{L} = the set of the vertices with heads. 1389
 1335 $\mathcal{R} = \{1, \dots, n\} - \mathcal{L}$ 1390
 1336 $\mathbf{B}_d^i = [\mathbf{A}_d]_{\mathcal{L}, \mathcal{R}}$ 1391
 1337 $\{\mathcal{X}^i, \mathcal{Y}^i\} = \arg \max_{|\mathcal{X}|+|\mathcal{Y}|=k} \mathbf{1}_{\mathcal{X}}^T \mathbf{B}_d^i \mathbf{1}_{\mathcal{Y}}$. 1392
 1338 **end for** 1393
 1339 $\{\mathcal{X}^i, \mathcal{Y}^i\} = \arg \max_{1 \leq i \leq n} \mathbf{1}_{\mathcal{X}^i}^T \mathbf{B}_d^i \mathbf{1}_{\mathcal{Y}^i}$ 1394
 1340 $\mathcal{Z}_d = \mathcal{X}_d \cup \mathcal{Y}_d$ 1395
 1341 21: **end if** 1396
 1342 22: **Output:** \mathcal{Z}_d 1397
 1343
 1344
 1345 **3.2. Proof of Theorem 2** 1399
 1346 We now proceed with with a simple bound that we get for the d -th largest eigenvalue of any adjacency matrix. 1400
 1347 This bound in combination with Theorem 1 will establish Theorem 2. 1401
 1348
 1349 **Proposition 10.** *For any un-weighted adjacency matrix \mathbf{A} , we have that* 1402
 1350
 1351
$$\lambda_d \leq \sqrt{\frac{2E}{d}}.$$
 1403
 1352
 1353 Moreover, for any un-weighted bi-adjacency matrix \mathbf{B} , we have that 1404
 1354
 1355
$$\sigma_d \leq \sqrt{\frac{E}{d}}.$$
 1405
 1356
 1357
 1358
 1359 *Proof.* First, observe that First, observe that 1410
 1360
 1361
$$d\lambda_d^2 \leq \sum_{i=1}^d \lambda_i^2 \leq \sum_{i=1}^n \lambda_i^2 = \|\mathbf{A}\|_F^2.$$
 1411
 1362
 1363
 1364 Observe moreover that 1412
 1365
$$\|\mathbf{A}\|_F = \sqrt{\sum_{i,j} A_{i,j}^2} = \sqrt{\sum_{i,j} A_{i,j}} = \sqrt{2 \cdot E},$$
 1413
 1366
 1367 where the second equality comes due to the fact that $A_{i,j}^2$ can only be 1 or 0. We now combine the above 1414
 1368 inequalities to obtain $\lambda_d \leq \sqrt{\frac{2E}{d}}$. For the rectangular case, we similarly have 1415
 1369
 1370
 1371
 1372
$$d\sigma_d^2 \leq \sum_{i=1}^d \sigma_i^2 \leq \sum_{i=1}^{\min(n_1, n_2)} \lambda_i^2 = \|\mathbf{B}\|_F^2.$$
 (17) 1416
 1373
 1374

$$\|\mathbf{A}\|_F = \sqrt{\sum_{i,j} A_{i,j}^2} = \sqrt{\sum_{i,j} A_{i,j}} = \sqrt{2 \cdot E},$$

where the second equality comes due to the fact that $A_{i,j}^2$ can only be 1 or 0. We now combine the above inequalities to obtain $\lambda_d \leq \sqrt{\frac{2E}{d}}$. For the rectangular case, we similarly have

$$d\sigma_d^2 \leq \sum_{i=1}^d \sigma_i^2 \leq \sum_{i=1}^{\min(n_1, n_2)} \lambda_i^2 = \|\mathbf{B}\|_F^2. \quad (17)$$

1430	Observe moreover that	1485
1431	$\ \mathbf{B}\ _F = \sqrt{\sum_{i,j} B_{i,j}^2} = \sqrt{\sum_{i,j} B_{i,j}} = \sqrt{E}$.	1486
1432		1487
1433		1488
1434	Hence, $\sigma_d \leq \sqrt{\frac{E}{d}}$. □	1489
1435		1490
1436	We can now establish our $1 + \epsilon$ result for bipartite matrices.	1491
1437	Claim 1. <i>If the densest-k-subgraph of \mathcal{G} contains a constant fraction of all the edges and \mathcal{G} is bipartite, and</i>	1492
1438	<i>$k = \Theta(\sqrt{E})$, then we can approximate DkS within a factor of $1 + \epsilon$, in time $n^{O(1/\epsilon^2)}$.</i>	1493
1439		1494
1440	<i>Proof.</i> Let	1495
1441		1496
1442	$\mathbf{A} = \begin{bmatrix} \mathbf{0}_{n_1 \times n_2} & \mathbf{B} \\ \mathbf{B}^T & \mathbf{0}_{n_2 \times n_1} \end{bmatrix},$	1497
1443		1498
1444	where $n_1 + n_2 = n$, and \mathbf{B} is the bi-adjacency matrix of the bipartite graph \mathcal{G} . Then, as we showed in Section 7	1499
1445	$\text{opt} = \max_{ \mathcal{S} =k} \mathbf{1}_S^T \mathbf{A} \mathbf{1}_S = 2 \max_{k_1, k_2: k=k_1+k_2} \max_{ \mathcal{X} =k_1} \max_{ \mathcal{Y} =k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}}$.	1500
1446		1501
1447	We approximate the above optimization, by solving the one on the rank- d matrix \mathbf{B}_d :	1502
1448		1503
1449	$\max_{k_1, k_2: k=k_1+k_2} \max_{ \mathcal{X} =k_1} \max_{ \mathcal{Y} =k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B}_d \mathbf{1}_{\mathcal{Y}}$.	1504
1450		1505
1451	and among all $O(k)$ pairs for all cases $k_1 + k_2 = k$, we output the one that is best. For each pair $k_1 \times k_2$, we	1506
1452	solve	1507
1453	$\max_{ \mathcal{X} =k_1} \max_{ \mathcal{Y} =k_2} \mathbf{1}_{\mathcal{X}}^T \mathbf{B}_d \mathbf{1}_{\mathcal{Y}}$.	1508
1454		1509
1455	Let $\{\mathcal{X}_d, \mathcal{Y}_d\}$ denote the solution, among all pairs of sparsity levels k_1, k_2 .	1510
1456	Then, due to Proposition 10, we have that	1511
1457		1512
1458	$\mathbf{1}_{\mathcal{X}_d}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}_d} \geq \mathbf{1}_{\mathcal{X}_*}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}_*} - 2 \cdot \sigma_{d+1}$.	1513
1459		1514
1460	We can now combine the spectral bound of Proposition 10 $\sigma_i \leq \sqrt{\frac{E}{i}}$ with our assumption that the optimal	1515
1461	k -subgraph contains a number of edges equal to $k \cdot \text{opt} = c \cdot E$, for some constant $c > 0$. Since $\sigma_d \leq \sqrt{\frac{E}{d}}$ and	1516
1462	$k = c\sqrt{E}$ we have	1517
1463		1518
1464	$2\text{opt}_d = 2\mathbf{1}_{\mathcal{X}_d}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}_d} \geq 2(\mathbf{1}_{\mathcal{X}_*}^T \mathbf{B} \mathbf{1}_{\mathcal{Y}_*} - 2 \cdot \sigma_{d+1})$ (18)	1519
1465		1520
1466	$\geq \text{opt} - 4 \cdot \sigma_{d+1} \geq c' \cdot \sqrt{E} - 4\sqrt{\frac{E}{d}} = \left(1 - 4\sqrt{\frac{2}{c^2 d}}\right) \text{opt}$ (19)	1521
1467		1522
1468		1523
1469	Hence, if we want $4\sqrt{\frac{2}{c^2 d}} = \epsilon$, we need	1524
1470	$d = \frac{16}{c^2 \epsilon^2} = O(1/\epsilon^2)$,	1525
1471		1526
1472	which establishes the result. □	1527
1473		1528
1474		1529
1475	Now, it is simple to establish our $2 + \epsilon$ result on general graphs that contain a dense subgraph.	1530
1476	Claim 2. <i>If the densest-k-subgraph of \mathcal{G} contains a constant fraction of all the edges and \mathcal{G}, and $k = \Theta(\sqrt{E})$, then we can approximate DkS within a factor of $2 + \epsilon$, in time $n^{O(1/\epsilon^2)}$.</i>	1531
1477		1532
1478		1533
1479	<i>Proof.</i> A simple combination of Proposition 10 and Proposition 4, give us our result. We first convert \mathcal{G} into a	1534
1480	bipartite with our random sampling Algorithm of Proposition 4. Then, we solve the bipartite problem on each	1535
1481	of the sampled graphs. The solution to the bipartite will be within $1 + \epsilon$, since the bipartite graph will contain	1536
1482	a constant fraction of the edges. Then converting it to a solution for the original graph, we lose an extra factor	1537
1483	of at most $2 + \delta$, as we showed in Proposition 4. □	1538
1484		1539

1540 3.3. Proof of Lemma 3: the data dependent bound

 1541 Finally, we prove our data dependent bounds, that bound the performance of opt, relative to polynomially
 1542 computable quantities.
 1543

 1544 **Claim 3.** *The optimal density of can be bounded as*

$$1545 \quad 1546 \quad \text{opt} \leq \min \left\{ \frac{2 + o_n(1)}{k} \cdot \mathbf{1}_{\mathcal{Z}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Z}_d} - |\lambda_{d+1}|, k - 1, \lambda_1 \right\}. \\ 1547$$

 1548 Additionally, if the graph is bipartite, or if the largest d eigenvalues of the graph are positive, then
 1549

$$1550 \quad 1551 \quad \text{opt} \leq \min \left\{ \frac{1}{k} \cdot \mathbf{1}_{\mathcal{Z}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Z}_d} - |\lambda_{d+1}|, k - 1, \lambda_1 \right\}. \\ 1552$$

 1553 *Proof.* Let $\mathcal{X}_*, \mathcal{Y}_*$ be the optimal solution of DBkS on \mathbf{A} and let $\mathcal{X}_d, \mathcal{Y}_d$ be the optimal solution of DBkS on
 1554 the rank- d matrix \mathbf{A}_d . Then, for the first bound we have
 1555

$$1556 \quad \text{opt}^B = \frac{\mathbf{1}_{\mathcal{X}_*}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}_*}}{k} = \frac{\mathbf{1}_{\mathcal{X}_*}^T (\mathbf{A}_d + \mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_*}}{k} = \frac{\mathbf{1}_{\mathcal{X}_*}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_*}}{k} + \frac{\mathbf{1}_{\mathcal{X}_*}^T (\mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_*}}{k} \\ 1557 \quad \leq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} + \frac{\mathbf{1}_{\mathcal{X}_*}^T (\mathbf{A} - \mathbf{A}_d) \mathbf{1}_{\mathcal{Y}_*}}{k} \leq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} + \frac{\|\mathbf{1}_{\mathcal{X}_*}\|_2 \|\mathbf{A} - \mathbf{A}_d\|_2 \|\mathbf{1}_{\mathcal{Y}_*}\|_2}{k} \leq \frac{\mathbf{1}_{\mathcal{X}_d}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}_d}}{k} + |\lambda_{d+1}|. \quad (20)$$

 1561 Applying this bound for all cases of graphs, as described by Theorem 1, and using the same bounds in the proofs
 1562 of Proposition 7-9, we obtain the result. The upper bound $k - 1$ is trivial by the fact that for any \mathcal{X} and \mathcal{Y} we
 1563 have

$$1564 \quad \frac{\mathbf{1}_{\mathcal{X}}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}}}{k} \leq \frac{\mathbf{1}_{\mathcal{X}}^T (\mathbf{1} \mathbf{1}^T - \mathbf{I}_n) \mathbf{1}_{\mathcal{Y}}}{k} \leq \frac{k(k-1)}{k} = k-1.$$

 1566 The last bound is simply due to the spectral bound on the bilinear form $\frac{\mathbf{1}_{\mathcal{X}}^T \mathbf{A} \mathbf{1}_{\mathcal{Y}}}{k} \leq \frac{\|\mathbf{1}_{\mathcal{X}}\|_2 \cdot \|\mathbf{A} \mathbf{1}_{\mathcal{Y}}\|_2}{k} \leq \lambda_1$. If we
 1567 combine now the above with the PSD part of Theorem 1, and the results of Proposition 4, we obtain the
 1568 claim. \square
 1569

 1570 4. Proof of Theorem 3: Nearly-linear Time Solver for Positive Semidefinite
 1571 Matrices

 1572 When the matrix \mathbf{A}_d has mixed signs of eigenvalues, then we have to go through the route of DBkS. However,
 1573 when \mathbf{A}_d has only positive eigenvalues, then it is easy to show that solving the bilinear problem on \mathbf{A}_d is
 1574 equivalent to solving

$$1577 \quad \max_{|\mathcal{X}|=k} \frac{\mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{X}}}{k}.$$

 1579 This is the DkS low-rank problem, that now can be solved by our algorithm. We show that when this spectral
 1580 scenario holds, we can speed up computations tremendously, by the use of a simple randomization.
 1581

 1582 Let us first remind the fact that DBkS and DkS are equivalent for positive semidefinite matrices. We will show
 1583 here how $\max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}}$ can be approximately in time nearly-linear in n , by only introducing a small relative
 1584 approximation error. Our approximation will use ϵ -nets.

 1585 **Definition 1.** (ϵ -net) Let $\mathbb{S}^d = \{\mathbf{c} \in \mathbb{R}^d : \|\mathbf{c}\|^2 = 1\}$ be the surface of the d -dimensional sphere. An ϵ -net of \mathbb{S}^d
 1586 is a finite set $\mathcal{N}_\epsilon^d \subset \mathbb{S}^d$ such that

$$1587 \quad \forall \mathbf{c} \in \mathbb{S}^d \ \exists \ \hat{\mathbf{c}} \in \mathcal{N}_\epsilon^d : \|\mathbf{c} - \hat{\mathbf{c}}\|_2 \leq \epsilon.$$

1588

 1589 We now show that we can solve our optimization, via the use of ϵ -nets.

 1590 **Proposition 11.** Let \mathcal{N}_ϵ^d be an ϵ -net of \mathbb{S}^d . Then,

$$1592 \quad (1 - \epsilon)^2 \cdot \max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}} \leq \max_{\mathbf{c} \in \mathcal{N}_\epsilon^d} \max_{|\mathcal{S}|=k} \left(\mathbf{c}^T \Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}} \right)^2 \leq \max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}}. \quad (21)$$

1650 Proof. Let \mathbf{c} be a $d \times 1$ unit length vector, i.e., $\|\mathbf{c}\|_2 = 1$. Then, by the Cauchy-Schwartz inequality we have
 1651
 1652 $(\mathbf{c}^T \Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}})^2 \leq \|\Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}}\|_2^2.$
 1653
 1654 We can get equality in the previous bound for a unit norm \mathbf{c} co-linear to $\Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}}$. Therefore, we have
 1655
 1656 $\|\Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}}\|_2^2 = \max_{\|\mathbf{c}\|_2=1} (\mathbf{c}^T \Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}})^2.$ (22)
 1657
 1658 Hence,
 1659
 1660 $\max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}} = \max_{|\mathcal{S}|=k} \|\Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}}\|_2^2 = \max_{|\mathcal{S}|=k} \max_{\|\mathbf{c}\|_2=1} (\mathbf{c}^T \Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}})^2.$ (23)
 1661
 1662 We can now obtain the upper bound of the proposition, since $\mathcal{N}_\epsilon^d \subseteq \mathbb{S}^d$. Now for the lower bound, let $(\mathbf{1}_{\mathcal{S}_d}, \mathbf{c}_d)$
 1663 denote the optimal solution of the above maximization, such that
 1664
 1665 $\max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}} = (\mathbf{c}_d^T \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}_d})^2.$
 1666
 1667 Then, there exists a vector $\hat{\mathbf{c}}$ in the ϵ -net \mathcal{N}_ϵ^d , such that $\mathbf{c}_d = \hat{\mathbf{c}} + \mathbf{r}$, with $\|\mathbf{r}\| \leq \epsilon$. Then,
 1668
 1669
 1670 $\sqrt{\max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}}} = \mathbf{c}_d^T \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}_d} = (\hat{\mathbf{c}} + \mathbf{r})^T \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}_d} \stackrel{(\alpha)}{\leq} \hat{\mathbf{c}}^T \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}_d} + \epsilon \cdot \|\mathbf{V}_d^T \mathbf{1}_{\mathcal{S}_d}\| = \hat{\mathbf{c}}^T \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}_d} + \epsilon \cdot \sqrt{\max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}}},$
 1671
 1672 where (α) is due to the triangle inequality, then the Cauchy-Schwartz inequality, and then the fact that $\|\mathbf{r}\| \leq \epsilon$.
 1673 From the above inequality we get
 1674
 1675
 1676 $(1 - \epsilon)^2 \max_{|\mathcal{S}|=k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{S}} \leq (\hat{\mathbf{c}}^T \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}_d})^2 \leq \max_{\mathbf{c} \in \mathcal{N}_\epsilon^d} \max_{|\mathcal{S}|=k} (\mathbf{c}^T \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}})^2$
 1677
 1678 which concludes the proof. \square
 1679
 1680 The importance of the above proposition is because for a fixed \mathbf{c} we can easily solve the problem
 1681
 1682
 1683 $\max_{|\mathcal{S}|=k} (\mathbf{c}^T \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}})^2.$
 1684
 1685 Observe that the above optimization is the same problem that we had to solve for a fixed \mathcal{Y} in Section 3. This
 1686 inner product is maximized when \mathcal{S} picks the largest, or smallest k elements of the n -dimensional vector $\mathbf{c}^T \mathbf{V}_d^T$.
 1687 The complexity to do that is linear $O(n)$ (Cormen et al., 2001).
 1688
 1689 It is now obvious that the number of elements, and the complexity to construct \mathcal{N}_ϵ^d is important. In (Wyner,
 1690 1967), it is shown how to randomly create an ϵ -net on the sphere, by randomly drawing m vectors on the sphere,
 1691 where
 1692 $m_\epsilon = \epsilon^{-d} \cdot e^{o(d)}.$
 1693
 1694 This approach gives with probability $1 - o_d(1)$ an ϵ -net. Let \mathcal{M}_ϵ be a set of m_ϵ vectors drawn uniformly on the
 1695 sphere. Our algorithm operates as follows: First we draw if we draw a set of \mathcal{M}_ϵ random vectors, and then we
 1696 find the corresponding optimal \mathcal{S} , by solving
 1697
 1698 $\max_{\mathbf{c} \in \mathcal{M}_\epsilon} \max_{|\mathcal{S}|=k} (\mathbf{c}^T \Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}})^2.$
 1699
 1700 This can be done in time $O(m_\epsilon n)$. Then, we draw the set \mathcal{M}_ϵ a total of $\log n$ times, and iterate this process
 1701 of finding the best \mathcal{S} for each of these sets \mathcal{M}_ϵ . Then, among all these solutions, with probability $1 - c/n$, for
 1702 some $c > 0$, the best solution satisfies the bound of the proposition. The randomized algorithm is given below
 1703 for completeness.
 1704

1760	Algorithm 6 Randomized Spannogram	1815
1761	1: <code>Spannogram_approx</code> (k , \mathbf{V}_d , Λ_d)	1816
1762	2: $\mathcal{S}_d = \emptyset$	1817
1763	3: for $i = 1 : m_\epsilon \cdot \log n$ do	1818
1764	4: $\mathbf{v} = (\Lambda_d^{1/2} \cdot \mathbf{V}_d)^T \cdot \text{randn}(d, 1)$	1819
1765	5: $\mathcal{S}_d = \mathcal{S}_d \cup \text{top}_k(\mathbf{v}) \cup \text{top}_k(-\mathbf{v})$	1820
1766	6: end for	1821
1767	7: Output: $\arg \max_{\mathcal{S} \in \mathcal{S}_d} \ \Lambda_d^{1/2} \mathbf{V}_d^T \mathbf{1}_{\mathcal{S}}\ .$	1822
1768		1823
1769	Computing \mathbf{A}_d in the first step of the algorithm, can also be done in nearly linear-time in the size of the input	1824
1770	\mathbf{A} . There is extensive literature on approximating \mathbf{A}_d by a rank- d matrix $\hat{\mathbf{A}}_d$ such that $\ \mathbf{A}_d - \hat{\mathbf{A}}_d\ _2 \leq \delta$.	1825
1771	in time proportional to the nonzero entries of \mathbf{A} times a poly-logarithmic term, as long as λ_d/λ_{d+1} is at least a	1826
1772	constant (Rokhlin et al., 2009; Halko et al., 2011; Gittens et al., 2013).	1827
1773		1828
1774	By further setting $\epsilon = \frac{1}{\log n}$, we obtain Theorem 3.	1829
1775		1830
1776	<h2>5. Vertex Sparsification via Simple Leverage Score Sampling</h2>	1831
1777		1832
1778	Our algorithm comes together with a vertex elimination step: after we compute the low-rank approximation	1833
1779	matrix \mathbf{A}_d , we discard rows and columns of \mathbf{A}_d , i.e., vertices, depending on their <i>weighted leverage scores</i> .	1834
1780	Leverage score sampling has been extensively studied in the literature, for many different applications, where	1835
1781	it can provably provide small error bounds, while keeping a small number of features from the original matrix	1836
1782	(Mahoney & Drineas, 2009; Boutsidis et al., 2009).	1837
1783		1838
1784	As we see in the following, this pre-processing step comes with an error guarantee. We show that by throwing	1839
1785	away vertices with small leverage scores, can only introduce a provably small error.	1840
1786	Let us define as	1841
1787	$\ell_i = \left\ \left[\mathbf{V}_d \Lambda ^{1/2} \right]_{i,:} \right\ = \sqrt{\sum_{j=1}^d [\mathbf{V}_d]_{i,j}^2 \lambda_j },$	1842
1788		1843
1789		1844
1790	the weighted leverage score of a vertex i . Then, our elimination step is simple. Let $\hat{\mathbf{A}}_d$ be a subset of \mathbf{A}_d , where	1845
1791	we have eliminated all vertices with $\ell_i \leq \frac{\eta}{3k\ell_1}$. Let $\mathbf{P}_{\mathcal{H}}$ be a diagonal matrix of 1s and 0s, with a 1 only in the	1846
1792	(i, i) indices such that $\ell_i > \frac{\eta}{3k\ell_1}$. Then,	1847
1793	$\hat{\mathbf{A}}_d = \mathbf{P}_{\mathcal{H}} \mathbf{A}_d \mathbf{P}_{\mathcal{H}}.$	1848
1794		1849
1795		1850
1796	We can now guarantee the following upper bound on the error introduced by the elimination.	1851
1797	Proposition 12. Let $\hat{\mathbf{A}}_d$ be created as above,	1852
1798		1853
1799	$\hat{\mathbf{A}}_d = \mathbf{P}_{\mathcal{H}} \mathbf{A}_d \mathbf{P}_{\mathcal{H}}.$	1854
1800	where $\mathbf{P}_{\mathcal{H}}$ is a diagonal matrix of 1s and 0s, with a 1 on (i, i) indices such that $\ell_i > \frac{\eta}{3k\ell_1}$. Then,	1855
1801		1856
1802	$\frac{\mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}}}{k} - \eta \leq \frac{\mathbf{1}_{\mathcal{X}}^T \hat{\mathbf{A}}_d \mathbf{1}_{\mathcal{Y}}}{k} \leq \frac{\mathbf{1}_{\mathcal{X}}^T \mathbf{A}_d \mathbf{1}_{\mathcal{Y}}}{k} + \eta,$	1857
1803		1858
1804		1859
1805	for all subsets of k vertices \mathcal{X}, \mathcal{Y} .	1860
1806		1861
1807	<i>Proof.</i> Let for brevity $\theta = \frac{\eta}{3k\ell_1}$. Moreover, Let $\mathbf{P}_{\mathcal{H}}$ be a diagonal matrix of 1s and 0s, with a 1 on (i, i) indices	1862
1808	such that $\ell_i > \frac{\eta}{3k\ell_1}$, and let $\mathbf{P}_{\mathcal{L}}$ be a diagonal matrix of 1s and 0s, with a 1 on (i, i) indices such that $\ell_i \leq \frac{\eta}{3k\ell_1}$.	1863
1809	Clearly,	1864
1810	$\mathbf{P}_{\mathcal{H}} + \mathbf{P}_{\mathcal{L}} = \mathbf{I}_{n \times n}.$	1865
1811		1866
1812	Then, we can rewrite \mathbf{A}_d as:	1867
1813	$\mathbf{A}_d = (\mathbf{P}_{\mathcal{H}} + \mathbf{P}_{\mathcal{L}}) \mathbf{A}_d (\mathbf{P}_{\mathcal{H}} + \mathbf{P}_{\mathcal{L}}) = \mathbf{P}_{\mathcal{H}} \mathbf{A}_d \mathbf{P}_{\mathcal{H}} + \mathbf{P}_{\mathcal{L}} \mathbf{A}_d \mathbf{P}_{\mathcal{L}} + \mathbf{P}_{\mathcal{H}} \mathbf{A}_d \mathbf{P}_{\mathcal{L}} + \mathbf{P}_{\mathcal{L}} \mathbf{A}_d \mathbf{P}_{\mathcal{H}}.$	1868
1814		1869

1870	Then, we have the following	1925
1871	$ \mathbf{1}_{\mathcal{X}}^T(\mathbf{A}_d - \hat{\mathbf{A}}_d)\mathbf{1}_y = \mathbf{1}_{\mathcal{X}}^T(\mathbf{P}_{\mathcal{L}}\mathbf{A}_d\mathbf{P}_{\mathcal{L}} + \mathbf{P}_{\mathcal{H}}\mathbf{A}_d\mathbf{P}_{\mathcal{L}} + \mathbf{P}_{\mathcal{L}}\mathbf{A}_d\mathbf{P}_{\mathcal{H}})\mathbf{1}_y $	1926
1872	$\leq \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{A}_d\mathbf{P}_{\mathcal{L}}\mathbf{1}_y + \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{H}}\mathbf{A}_d\mathbf{P}_{\mathcal{L}}\mathbf{1}_y + 2 \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{A}_d\mathbf{P}_{\mathcal{H}}y $	1927
1873		1928
1874		1929
1875	Observe that for the first error term we have	1930
1876		1931
1877	$ \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{A}_d\mathbf{P}_{\mathcal{L}}\mathbf{1}_y = \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{V}_d\Lambda_d^{1/2}\mathbf{S}\Lambda_d^{1/2}\mathbf{V}_d^T\mathbf{P}_{\mathcal{L}}\mathbf{1}_y $	1932
1878	$\leq \ \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{V}_d\Lambda_d^{1/2}\ \cdot \ \mathbf{S}\Lambda_d^{1/2}\mathbf{V}_d^T\mathbf{P}_{\mathcal{L}}\mathbf{1}_y\ \leq \ \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{V}_d\Lambda_d^{1/2}\ \cdot \ \mathbf{S}\ \cdot \ \Lambda_d^{1/2}\mathbf{V}_d^T\mathbf{P}_{\mathcal{L}}\mathbf{1}_y\ $	1933
1879		1934
1880	$= \ \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{V}_d\Lambda_d^{1/2}\ \cdot 1 \cdot \ \Lambda_d^{1/2}\mathbf{V}_d^T\mathbf{P}_{\mathcal{L}}\mathbf{1}_y\ \leq k \cdot \theta \cdot k \cdot \theta = k^2 \cdot \theta^2.$	1935
1881	where the second and third inequalities come due to the Cauchy-Schwarz inequality. In the above \mathbf{S} denotes the diagonal matrix that contains the signs of the eigenvalues. Clearly, its operator norm is 1. Hence, the last	1936
1882	inequality in the above is due to the fact that $\mathbf{1}_{\mathcal{X}}, \mathbf{1}_y$ have k entries with 1, and each picks the rows of $\mathbf{V}_d\Lambda_d^{1/2}$	1937
1883	with the highest leverage score. Then, due to the triangle inequality on the k -largest row norms (i.e., leverage	1938
1884	scores) of $\mathbf{V}_d\Lambda_d^{1/2}$ we get the final result.	1939
1885		1940
1886		1941
1887	Similarly, we can bound the remaining two error terms	1942
1888		1943
1889	$ \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{A}_d\mathbf{P}_{\mathcal{H}}\mathbf{1}_y = \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{V}_d\Lambda_d^{1/2}\mathbf{S}\Lambda_d^{1/2}\mathbf{V}_d^T\mathbf{P}_{\mathcal{H}}\mathbf{1}_y $	1944
1890	$\leq \ \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{V}_d\Lambda_d^{1/2}\ \cdot \ \mathbf{S}\Lambda_d^{1/2}\mathbf{V}_d^T\mathbf{P}_{\mathcal{H}}\mathbf{1}_y\ \leq \ \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{V}_d\Lambda_d^{1/2}\ \cdot \ \mathbf{S}\ \cdot \ \Lambda_d^{1/2}\mathbf{V}_d^T\mathbf{P}_{\mathcal{H}}\mathbf{1}_y\ $	1945
1891		1946
1892	$= \ \mathbf{1}_{\mathcal{X}}^T\mathbf{P}_{\mathcal{L}}\mathbf{V}_d\Lambda_d^{1/2}\ \cdot 1 \cdot \ \Lambda_d^{1/2}\mathbf{V}_d^T\mathbf{P}_{\mathcal{H}}\mathbf{1}_y\ \leq k \cdot \theta \cdot k \cdot \ell_1 = k^2 \cdot \theta \cdot \ell_1.$	1947
1893	Since, $\eta \leq \ell_1$, we conclude that the above error can be bounded as	1948
1894		1949
1895	$\frac{ \mathbf{1}_{\mathcal{X}}^T(\mathbf{A}_d - \hat{\mathbf{A}}_d)\mathbf{1}_y }{k} \leq 3 \cdot k \cdot \theta \cdot \ell_1 = \eta.$	1950
1896		1951
1897	Hence, we obtain the proposition. \square	1952
1898		1953
1899		1954
1900	6. NP-hardness of DkS on rank-1 matrices	1955
1901	In this section, we establish the hardness of the quadratic formulation of DkS, even for rank-1 matrices. Interestingly the problem is not hard when we relax it to its bilinear form as we showed in our main result. The claim follows.	1956
1902		1957
1903		1958
1904		1959
1905	Claim 4. <i>DkS is NP-hard for rank-1 matrices \mathbf{A} with one negative eigenvalue.</i>	1960
1906		1961
1907	<i>Proof.</i> Observe that a rank-1 matrix with 1 negative eigenvalue can be written as	1962
1908		1963
1909	$\mathbf{A} = -\mathbf{v}\mathbf{v}^T$	1964
1910	where $\lambda_1 = \ \mathbf{v}\ ^2$. Then, see that	1965
1911		1966
1912	$\max_{ \mathcal{S} =k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A} \mathbf{1}_{\mathcal{S}} = \max_{ \mathcal{S} =k} -\mathbf{1}_{\mathcal{S}}^T \mathbf{v} \mathbf{v}^T \mathbf{1}_{\mathcal{S}} = \min_{ \mathcal{S} =k} \mathbf{1}_{\mathcal{S}}^T \mathbf{v} \mathbf{v}^T \mathbf{1}_{\mathcal{S}} = \left(\min_{ \mathcal{S} =k} \mathbf{1}_{\mathcal{S}}^T \mathbf{v} \right)^2 = \left(\min_{ \mathcal{S} =k} \left \sum_{i \in \mathcal{S}} v_i \right \right)^2.$	1967
1913		1968
1914	An algorithm that can solve the above problem, can be used to solve SUBSETSUM. In SUBSETSUM we are given	1969
1915	a set of integers and we wish to decide whether there exists a non-empty subset of these integers that sums to	1970
1916	zero. In the following algorithm we show how this can be trivially done, by solving $\min_{ \mathcal{S} =k} \sum_{i \in \mathcal{S}} v_i $ for all	1971
1917	values of k . If for some value of k the sum in the optimizaton is zero, then we decide YES as the output for the	1972
1918	SUBSETSUM. Hence, solving $\max_{ \mathcal{S} =k} \mathbf{1}_{\mathcal{S}}^T \mathbf{A} \mathbf{1}_{\mathcal{S}}$ is NP-hard even for rank-1 matrices, in the general case. \square	1973
1919		1974
1920	7. Additional Experiments	1975
1921		1976
1922	In Fig. 1, we show additional experiment on 9 more large-graphs. The description of the graphs can be found in	1977
1923	Table 1. Moreover, in Fig. 2. The description of the experiments can be found in the figure captions.	1978
1924		1979

1980	Algorithm 7 SubsetSum via rank-1 DkS	2035			
1981	1: Input: $v = [v_1, \dots, v_n]$	2036			
1982	2: for $i = 1 : k$ do	2037			
1983	3: $s_i = \min_{ \mathcal{S} =k} \sum_{i \in \mathcal{S}} v_i $	2038			
1984	4: if $s_i == 0$ then	2039			
1985	5: Output: YES	2040			
1986	6: else	2041			
1987	7: Output: NO	2042			
1988	8: end if	2043			
1989	9: end for	2044			
1990		2045			
1991		2046			
1992	<i>Table 1.</i> Datasets used in our experiments				
1993		2047			
1994		2048			
1995	DATA SET	NODES	EDGES	DESCRIPTION	2049
1996					2050
1997	COM-DBLP	317,080	1,049,866	DBLP COLLABORATION NETWORK	2051
1998	COM-LIVEJOURNAL	3,997,962	34,681,189	LIVEJOURNAL ONLINE SOCIAL NETWORK	2052
1999	WEB-NOTREDAME	325,729	1,497,134	WEB GRAPH OF NOTRE DAME	2053
2000	EGO-FACEBOOK	4,039	88,234	SOCIAL CIRCLES FROM FACEBOOK (ANONYMIZED)	2054
2001	CA-ASTROPH	18,772	396,160	COLLABORATION NETWORK OF ARXIV ASTRO PHYSICS	2055
2002	CA-HEPPH	12,008	237,010	COLLABORATION NETWORK OF ARXIV HIGH ENERGY PHYSICS	2056
2003	CA-CONDMAT	23,133	186,936	COLLABORATION NETWORK OF ARXIV CONDENSED MATTER	2057
2004	CA-GRQC	5,242	28,980	COLLABORATION NETWORK OF ARXIV GENERAL RELATIVITY	2058
2005	CA-HEPTH	9,877	51,971	COLLABORATION NETWORK OF ARXIV HIGH ENERGY PHYSICS THEORY	2059
2006	LOC-BRIGHTKITE	58,228	214,078	BRIGHTKITE LOCATION BASED ONLINE SOCIAL NETWORK	2060
2007	ROADNET-CA	1,965,206	5,533,214	ROAD NETWORK OF CALIFORNIA	2061
2008	EMAIL-ENRON	36,692	367,662	EMAIL COMMUNICATION NETWORK FROM ENRON	2062
2009	COM-ORKUT	3,072,441	117,185,083	ORKUT ONLINE SOCIAL NETWORK	2063
2010	FLICKR	105,938	2,316,948	NETWORK OF FLICKR IMAGES SHARING COMMON METADATA	2064
2011	FBMEDIUM	63,731	817,090	FRIENDSHIP DATA OF FACEBOOK USERS	2065

References

- 2013 Boutsidis, Christos, Mahoney, Michael W, and Drineas, Petros. An improved approximation algorithm for the column subset selection problem. In *Proceedings of the twentieth Annual ACM-SIAM Symposium on Discrete Algorithms*, pp. 968–977. Society for Industrial and Applied Mathematics, 2009.
- 2014 Cormen, Thomas H, Leiserson, Charles E, Rivest, Ronald L, and Stein, Clifford. *Introduction to algorithms*. MIT press, 2001.
- 2015 Gittens, Alex, Kambadur, Prabhanjan, and Boutsidis, Christos. Approximate spectral clustering via randomized sketching. *arXiv preprint arXiv:1311.2854*, 2013.
- 2016 Halko, Nathan, Martinsson, Per-Gunnar, and Tropp, Joel A. Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions. *SIAM review*, 53(2):217–288, 2011.
- 2017 Mahoney, Michael W and Drineas, Petros. Cur matrix decompositions for improved data analysis. *Proceedings of the National Academy of Sciences*, 106(3):697–702, 2009.
- 2018 Papailiopoulos, Dimitris S, Dimakis, Alexandros G, and Korokithakis, Stavros. Sparse pca through low-rank approximations. *arXiv preprint arXiv:1303.0551*, 2013.
- 2019 Rokhlin, Vladimir, Szlam, Arthur, and Tygert, Mark. A randomized algorithm for principal component analysis. *SIAM Journal on Matrix Analysis and Applications*, 31(3):1100–1124, 2009.
- 2020 Wyner, Aaron D. Random packings and coverings of the unit n-sphere. *Bell System Technical Journal*, 46(9):2111–2118, 1967.

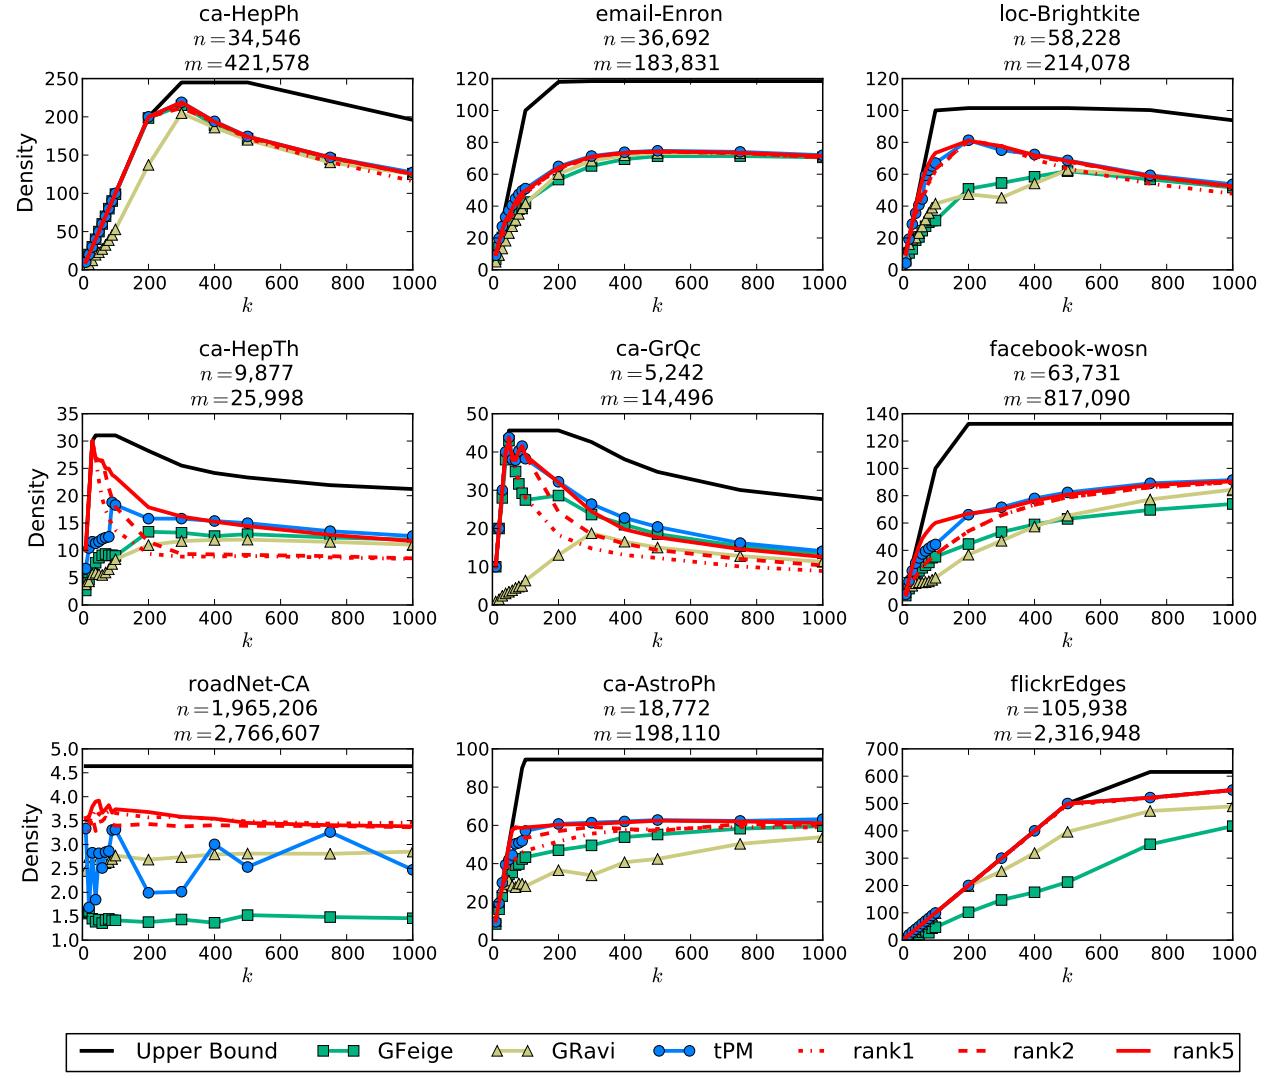


Figure 1. Subgraph density vs. subgraph size (k). We show the comparison of densest subgraph algorithms on several additional datasets: Academic collaboration graphs from Arxiv (ca-HepTh, ca-HepPh, ca-GrQc, Ca-Astro), Geographic location-based networks (roadNet, loc-Brightkite), The Enron email communication graph (email-Enron) and a facebook subgraph (facebook-wosn). The number of vertices and edges are shown in each plot. As can be seen, in almost all cases rank-2 and rank-5 spannograms match or outperform previous algorithms. One notable exception is the ca-GrQc where, for subgraphs of size above $k = 400$ or above, T-power performs better. Another observation is that the tightness of our data-dependent bound (solid black line) varies for different data sets and subgraph sizes.

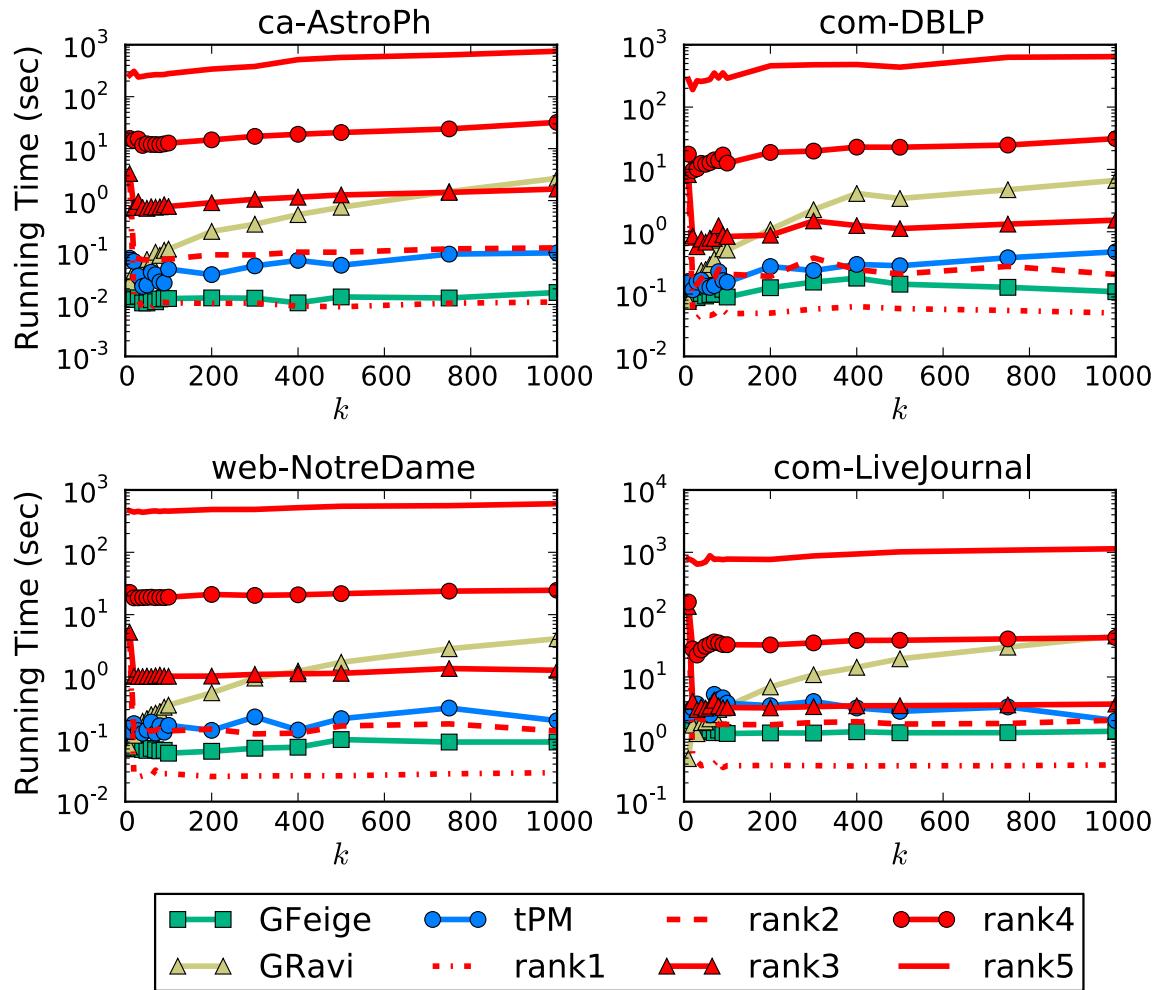


Figure 2. Running times on a MacBook Pro 10.2, with Intel Core i5 @ 2.5 GHz (2 cores), 256 KB L2 Cache per core, 3 MB L3 Cache, and 8 GB RAM. Experiments were run on MATLAB R2011b (7.13.0.564). As can be seen, Rank-1 is significantly faster than all other algorithms for all tested cases. Rank-2 is comparable to prior work, having running times of a few seconds. Rank-5 was the highest accuracy setting we tested. It can take several minutes on large graphs and seems useful only when high accuracy is desired or other methods are far from the upper bound.